

# Exercises

February 21, 2025

```
[1]: from notebookfuncs import *
```

```
[2]: import numpy as np
import matplotlib.pyplot as plt
from matplotlib.pyplot import subplots
import pandas as pd
from ISLP import load_data
import seaborn as sns
from numpy import median
```

```
[3]: College = pd.read_csv("College.csv")
College
```

```
[3]:
```

	Unnamed: 0	Private	Apps	Accept	Enroll	Top10perc	\
0	Abilene Christian University	Yes	1660	1232	721	23	
1	Adelphi University	Yes	2186	1924	512	16	
2	Adrian College	Yes	1428	1097	336	22	
3	Agnes Scott College	Yes	417	349	137	60	
4	Alaska Pacific University	Yes	193	146	55	16	
..	...	...	...	...	...	...	
772	Worcester State College	No	2197	1515	543	4	
773	Xavier University	Yes	1959	1805	695	24	
774	Xavier University of Louisiana	Yes	2097	1915	695	34	
775	Yale University	Yes	10705	2453	1317	95	
776	York College of Pennsylvania	Yes	2989	1855	691	28	

	Top25perc	F.Undergrad	P.Undergrad	Outstate	Room.Board	Books	\
0	52	2885	537	7440	3300	450	
1	29	2683	1227	12280	6450	750	
2	50	1036	99	11250	3750	400	
3	89	510	63	12960	5450	450	
4	44	249	869	7560	4120	800	
..	...	...	...	...	...	...	
772	26	3089	2029	6797	3900	500	
773	47	2849	1107	11520	4960	600	
774	61	2793	166	6900	4200	617	
775	99	5217	83	19840	6510	630	
776	63	2988	1726	4990	3560	500	

	Personal	PhD	Terminal	S.F.Ratio	perc.alumni	Expend	Grad.Rate
0	2200	70	78	18.1	12	7041	60
1	1500	29	30	12.2	16	10527	56
2	1165	53	66	12.9	30	8735	54
3	875	92	97	7.7	37	19016	59
4	1500	76	72	11.9	2	10922	15
..	...	...	...	...	...	...	...
772	1200	60	60	21.0	14	4469	40
773	1250	73	75	13.3	31	9189	83
774	781	67	75	14.4	20	8323	49
775	2115	96	96	5.8	49	40386	99
776	1250	75	75	18.1	28	4509	99

[777 rows x 19 columns]

```
[4]: college2 = pd.read_csv("College.csv", index_col=0)
college2
```

```
[4]:
```

	Private	Apps	Accept	Enroll	Top10perc	\
Abilene Christian University	Yes	1660	1232	721	23	
Adelphi University	Yes	2186	1924	512	16	
Adrian College	Yes	1428	1097	336	22	
Agnes Scott College	Yes	417	349	137	60	
Alaska Pacific University	Yes	193	146	55	16	
...	...	...	...	...	...	
Worcester State College	No	2197	1515	543	4	
Xavier University	Yes	1959	1805	695	24	
Xavier University of Louisiana	Yes	2097	1915	695	34	
Yale University	Yes	10705	2453	1317	95	
York College of Pennsylvania	Yes	2989	1855	691	28	

	Top25perc	F.Undergrad	P.Undergrad	Outstate	\
Abilene Christian University	52	2885	537	7440	
Adelphi University	29	2683	1227	12280	
Adrian College	50	1036	99	11250	
Agnes Scott College	89	510	63	12960	
Alaska Pacific University	44	249	869	7560	
...	...	...	...	...	
Worcester State College	26	3089	2029	6797	
Xavier University	47	2849	1107	11520	
Xavier University of Louisiana	61	2793	166	6900	
Yale University	99	5217	83	19840	
York College of Pennsylvania	63	2988	1726	4990	

	Room.Board	Books	Personal	PhD	Terminal	\
Abilene Christian University	3300	450	2200	70	78	

Adelphi University	6450	750	1500	29	30
Adrian College	3750	400	1165	53	66
Agnes Scott College	5450	450	875	92	97
Alaska Pacific University	4120	800	1500	76	72
...	...	...	...	...	...
Worcester State College	3900	500	1200	60	60
Xavier University	4960	600	1250	73	75
Xavier University of Louisiana	4200	617	781	67	75
Yale University	6510	630	2115	96	96
York College of Pennsylvania	3560	500	1250	75	75

	S.F.Ratio	perc.alumni	Expend	Grad.Rate
Abilene Christian University	18.1	12	7041	60
Adelphi University	12.2	16	10527	56
Adrian College	12.9	30	8735	54
Agnes Scott College	7.7	37	19016	59
Alaska Pacific University	11.9	2	10922	15
...	...	...	...	...
Worcester State College	21.0	14	4469	40
Xavier University	13.3	31	9189	83
Xavier University of Louisiana	14.4	20	8323	49
Yale University	5.8	49	40386	99
York College of Pennsylvania	18.1	28	4509	99

[777 rows x 18 columns]

```
[5]: College3 = College.rename({"Unnamed: 0": "College"}, axis=1)
College3.set_index("College")
College3
```

```
[5]:
```

	College	Private	Apps	Accept	Enroll	Top10perc	\
0	Abilene Christian University	Yes	1660	1232	721	23	
1	Adelphi University	Yes	2186	1924	512	16	
2	Adrian College	Yes	1428	1097	336	22	
3	Agnes Scott College	Yes	417	349	137	60	
4	Alaska Pacific University	Yes	193	146	55	16	
..	...	...	...	...	...	...	
772	Worcester State College	No	2197	1515	543	4	
773	Xavier University	Yes	1959	1805	695	24	
774	Xavier University of Louisiana	Yes	2097	1915	695	34	
775	Yale University	Yes	10705	2453	1317	95	
776	York College of Pennsylvania	Yes	2989	1855	691	28	

	Top25perc	F.Undergrad	P.Undergrad	Outstate	Room.Board	Books	\
0	52	2885	537	7440	3300	450	
1	29	2683	1227	12280	6450	750	
2	50	1036	99	11250	3750	400	

3	89	510	63	12960	5450	450
4	44	249	869	7560	4120	800
..	...	...	...	...	...	...
772	26	3089	2029	6797	3900	500
773	47	2849	1107	11520	4960	600
774	61	2793	166	6900	4200	617
775	99	5217	83	19840	6510	630
776	63	2988	1726	4990	3560	500

	Personal	PhD	Terminal	S.F.Ratio	perc.alumni	Expend	Grad.Rate
0	2200	70	78	18.1	12	7041	60
1	1500	29	30	12.2	16	10527	56
2	1165	53	66	12.9	30	8735	54
3	875	92	97	7.7	37	19016	59
4	1500	76	72	11.9	2	10922	15
..	...	...	...	...	...	...	...
772	1200	60	60	21.0	14	4469	40
773	1250	73	75	13.3	31	9189	83
774	781	67	75	14.4	20	8323	49
775	2115	96	96	5.8	49	40386	99
776	1250	75	75	18.1	28	4509	99

[777 rows x 19 columns]

[6]: College = College3

[6]:

	College	Private	Apps	Accept	Enroll	Top10perc	\
0	Abilene Christian University	Yes	1660	1232	721	23	
1	Adelphi University	Yes	2186	1924	512	16	
2	Adrian College	Yes	1428	1097	336	22	
3	Agnes Scott College	Yes	417	349	137	60	
4	Alaska Pacific University	Yes	193	146	55	16	
..	...	...	...	...	...	...	...
772	Worcester State College	No	2197	1515	543	4	
773	Xavier University	Yes	1959	1805	695	24	
774	Xavier University of Louisiana	Yes	2097	1915	695	34	
775	Yale University	Yes	10705	2453	1317	95	
776	York College of Pennsylvania	Yes	2989	1855	691	28	

	Top25perc	F.Undergrad	P.Undergrad	Outstate	Room.Board	Books	\
0	52	2885	537	7440	3300	450	
1	29	2683	1227	12280	6450	750	
2	50	1036	99	11250	3750	400	
3	89	510	63	12960	5450	450	
4	44	249	869	7560	4120	800	
..	...	...	...	...	...	...	...
772	26	3089	2029	6797	3900	500	

773	47	2849	1107	11520	4960	600
774	61	2793	166	6900	4200	617
775	99	5217	83	19840	6510	630
776	63	2988	1726	4990	3560	500

	Personal	PhD	Terminal	S.F.Ratio	perc.alumni	Expend	Grad.Rate
0	2200	70	78	18.1	12	7041	60
1	1500	29	30	12.2	16	10527	56
2	1165	53	66	12.9	30	8735	54
3	875	92	97	7.7	37	19016	59
4	1500	76	72	11.9	2	10922	15
..	...	...	...	...	...	...	...
772	1200	60	60	21.0	14	4469	40
773	1250	73	75	13.3	31	9189	83
774	781	67	75	14.4	20	8323	49
775	2115	96	96	5.8	49	40386	99
776	1250	75	75	18.1	28	4509	99

[777 rows x 19 columns]

```
[7]: College.describe()
```

```
[7]:
```

	Apps	Accept	Enroll	Top10perc	Top25perc	\
count	777.000000	777.000000	777.000000	777.000000	777.000000	
mean	3001.638353	2018.804376	779.972973	27.558559	55.796654	
std	3870.201484	2451.113971	929.176190	17.640364	19.804778	
min	81.000000	72.000000	35.000000	1.000000	9.000000	
25%	776.000000	604.000000	242.000000	15.000000	41.000000	
50%	1558.000000	1110.000000	434.000000	23.000000	54.000000	
75%	3624.000000	2424.000000	902.000000	35.000000	69.000000	
max	48094.000000	26330.000000	6392.000000	96.000000	100.000000	

	F.Undergrad	P.Undergrad	Outstate	Room.Board	Books	\
count	777.000000	777.000000	777.000000	777.000000	777.000000	
mean	3699.907336	855.298584	10440.669241	4357.526384	549.380952	
std	4850.420531	1522.431887	4023.016484	1096.696416	165.105360	
min	139.000000	1.000000	2340.000000	1780.000000	96.000000	
25%	992.000000	95.000000	7320.000000	3597.000000	470.000000	
50%	1707.000000	353.000000	9990.000000	4200.000000	500.000000	
75%	4005.000000	967.000000	12925.000000	5050.000000	600.000000	
max	31643.000000	21836.000000	21700.000000	8124.000000	2340.000000	


  

	Personal	PhD	Terminal	S.F.Ratio	perc.alumni	\
count	777.000000	777.000000	777.000000	777.000000	777.000000	
mean	1340.642214	72.660232	79.702703	14.089704	22.743887	
std	677.071454	16.328155	14.722359	3.958349	12.391801	
min	250.000000	8.000000	24.000000	2.500000	0.000000	

25%	850.000000	62.000000	71.000000	11.500000	13.000000
50%	1200.000000	75.000000	82.000000	13.600000	21.000000
75%	1700.000000	85.000000	92.000000	16.500000	31.000000
max	6800.000000	103.000000	100.000000	39.800000	64.000000


	Expend	Grad.Rate
count	777.000000	777.000000
mean	9660.171171	65.46332
std	5221.768440	17.17771
min	3186.000000	10.00000
25%	6751.000000	53.00000
50%	8377.000000	65.00000
75%	10830.000000	78.00000
max	56233.000000	118.00000

```
[8]: pd.plotting.scatter_matrix(College[["Top10perc", "Apps", "Enroll"]]);
```



```
[9]: fig, ax = subplots(figsize=(8, 8))
College.boxplot("Outstate", by="Private", ax=ax);
```

```
Executing <Handle IOLoop._run_callback(functools.par...7dd256f2f600>)) created
at /home/linux/ISLP/islpenv/lib/python3.12/site-
packages/tornado/platform/asyncio.py:235> took 0.122 seconds
IOStream.flush timed out
Executing <Handle BaseAsyncIOLoop._handle_events(28, 1) created at
/usr/lib/python3.12/asyncio/selector_events.py:280> took 0.134 seconds
```



```
[10]: College["Top10perc"]
```

```
[10]: 0      23
      1      16
      2      22
```

```
3      60
4      16
      ..
772     4
773    24
774    34
775    95
776    28
```

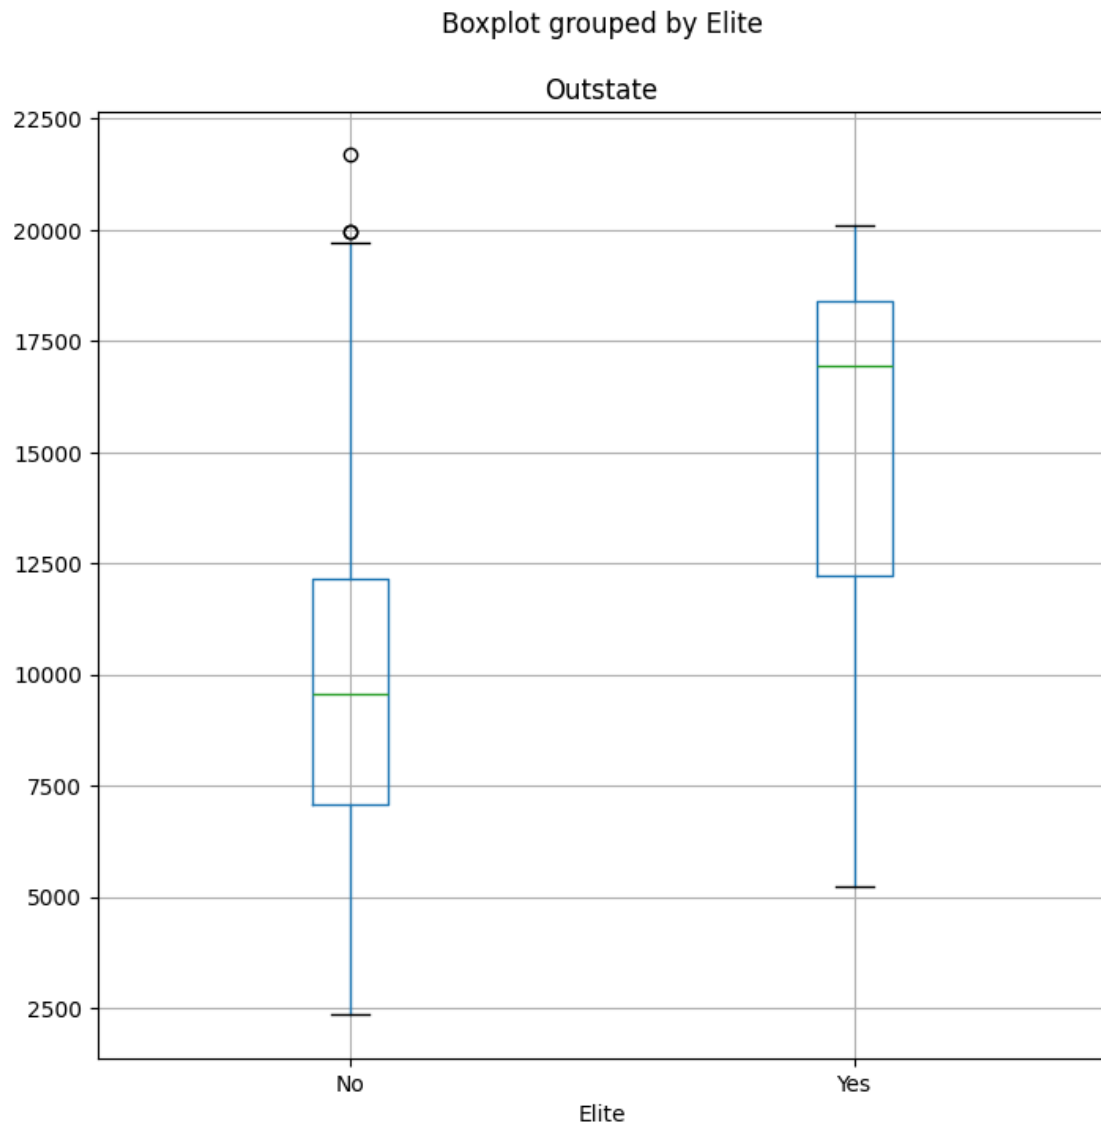
Name: Top10perc, Length: 777, dtype: int64

```
[11]: College["Elite"] = pd.cut(College["Top10perc"], [0, 50, 100], labels=["No", "Yes"])
      College["Elite"].value_counts()
```


```
[11]: Elite
      No      699
      Yes      78
      Name: count, dtype: int64
```

```
[12]: fig, ax = subplots(figsize=(8, 8))
      College.boxplot("Outstate", by="Elite", ax=ax);
```

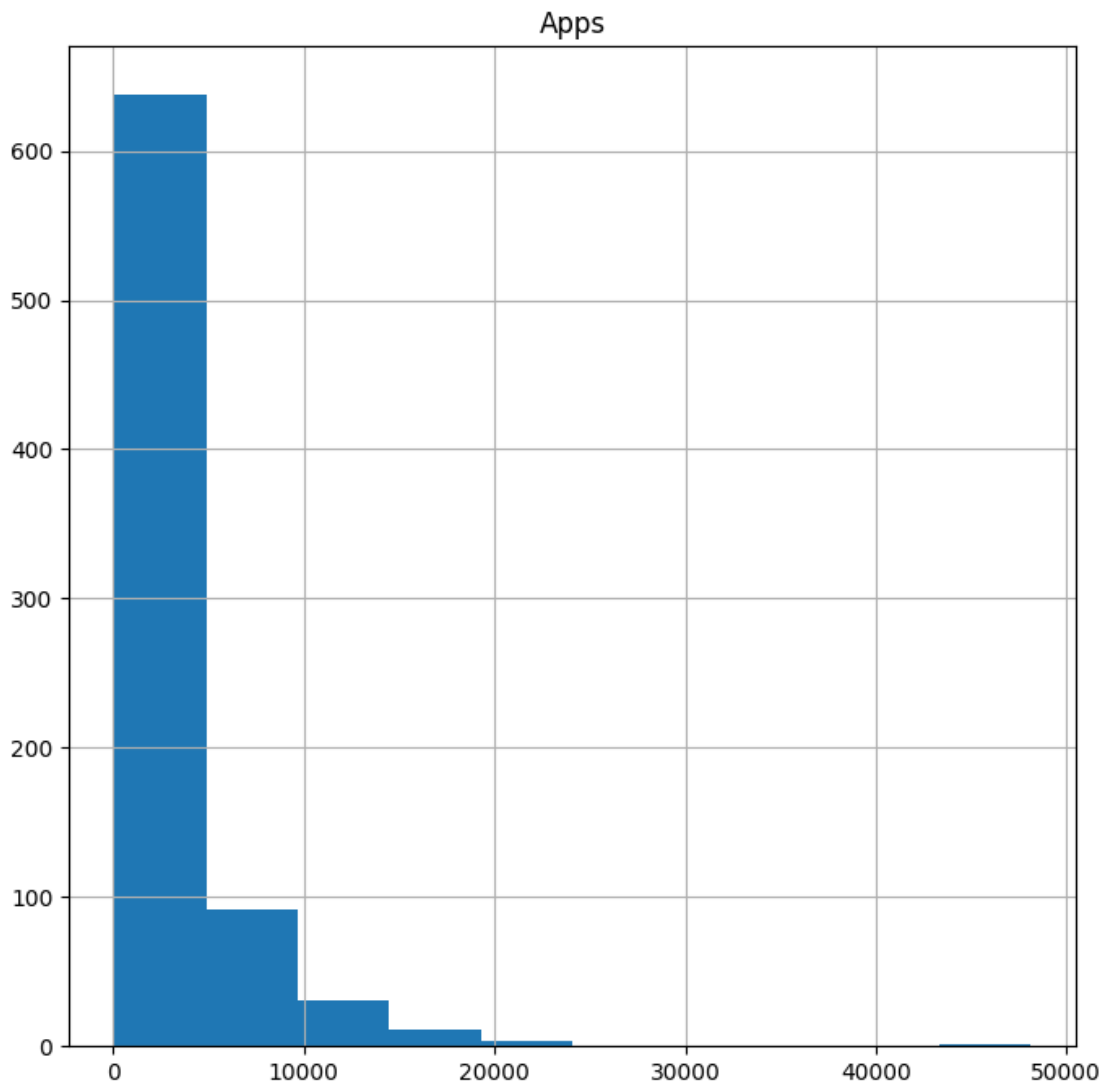




```
[13]: College.plot.hist();
```



```
[14]: fig, ax = subplots(figsize=(8, 8))
      College.hist("Apps", ax=ax);
```



```
[15]: numeric_columns = College.select_dtypes(include="number").columns.tolist()
      numeric_columns
```

```
[15]: ['Apps',
      'Accept',
      'Enroll',
      'Top10perc',
      'Top25perc',
      'F.Undergrad',
      'P.Undergrad',
      'Outstate',
      'Room.Board',
      'Books',
```

```

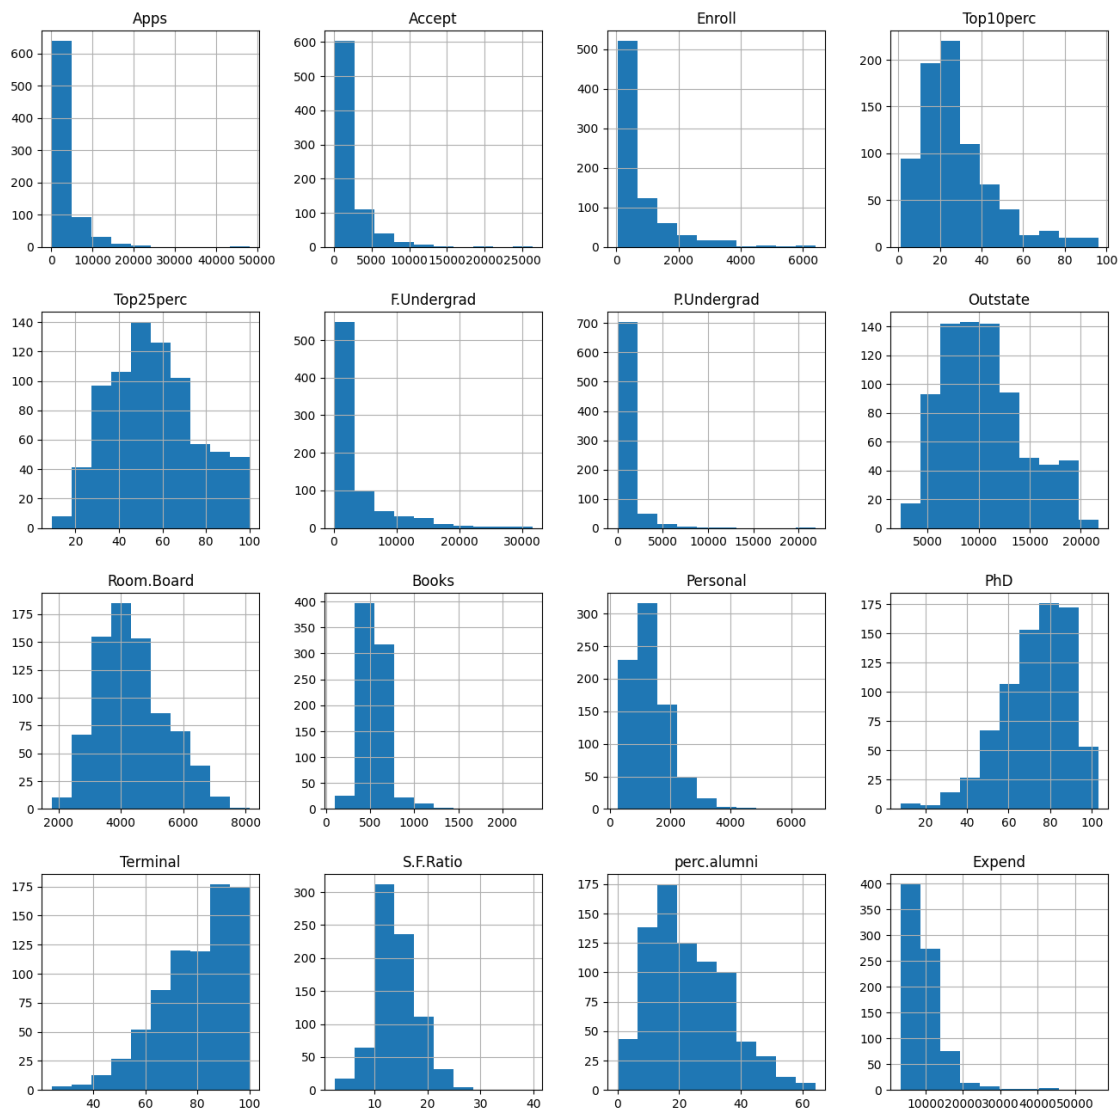
'Personal',
'PhD',
'Terminal',
'S.F.Ratio',
'perc.alumni',
'Expend',
'Grad.Rate']

```

```

[16]: fig, axs = subplots(4, 4, figsize=(16, 16))
      for row in range(0, 4):
          for column in range(0, 4):
              College.hist(numeric_columns[row * 4 + column], ax=axs[row, column])

```



### 0.0.1 Count of private and public colleges

```
[17]: College["Private"].value_counts()
```


```
[17]: Private  
Yes    565  
No     212  
Name: count, dtype: int64
```

```
[18]: College["AcceptanceRate"] = round(College["Accept"] / College["Apps"] * 100, 2)  
College["AcceptanceRate"]
```

```
[18]: 0      74.22  
     1      88.01  
     2      76.82  
     3      83.69  
     4      75.65  
     ...  
    772      68.96  
    773      92.14  
    774      91.32  
    775      22.91  
    776      62.06  
     Name: AcceptanceRate, Length: 777, dtype: float64
```


```
[19]: ### Plot boxplot for acceptance rate by College Type : Elite or not
```

```
[20]: fig, ax = subplots(figsize=(8, 8))  
      College.boxplot("AcceptanceRate", by="Elite", ax=ax);
```



```
[21]: ### Plot boxplot for acceptance rate for Private colleges or not
```

```
[22]: fig, ax = subplots(figsize=(8, 8))  
College.boxplot("AcceptanceRate", by="Private", ax=ax);
```




```
[23]: College["EnrollmentRate"] = round(College["Enroll"] / College["Accept"] * 100, 2)
      College["EnrollmentRate"]
```

```
[23]: 0      58.52
      1      26.61
      2      30.63
      3      39.26
      4      37.67
      ...
      772    35.84
```


```
773    38.50
774    36.29
775    53.69
776    37.25
Name: EnrollmentRate, Length: 777, dtype: float64
```

```
[24]: fig, ax = subplots(figsize=(8, 8))
College.boxplot("EnrollmentRate", by="Elite", ax=ax);
```




```
[25]: fig, ax = subplots(figsize=(8, 8))
College.boxplot("EnrollmentRate", by="Private", ax=ax);
```






```
[26]: College["StudentCosts"] = (
    College["Outstate"] + College["Room.Board"] + College["Books"] +
    ↪ College["Personal"]
)
```


```
[27]: fig, ax = subplots(figsize=(8, 8))
College.boxplot("StudentCosts", by="Elite", ax=ax);
```




```
[28]: fig, ax = subplots(figsize=(8, 8))  
College.boxplot("StudentCosts", by="Private", ax=ax);
```




```
[29]: fig, ax = subplots(figsize=(8, 8))  
College.boxplot("PhD", by="Private", ax=ax);
```




```
[30]: fig, ax = subplots(figsize=(8, 8))  
College.boxplot("PhD", by="Elite", ax=ax);
```




```
[31]: fig, ax = subplots(figsize=(8, 8))  
College.boxplot("Terminal", by="Elite", ax=ax);
```



```
[32]: fig, ax = subplots(figsize=(8, 8))  
College.boxplot("Terminal", by="Private", ax=ax);
```




```
[33]: fig, ax = subplots(figsize=(8, 8))  
College.boxplot("S.F.Ratio", by="Private", ax=ax);
```




```
[34]: fig, ax = subplots(figsize=(8, 8))  
College.boxplot("S.F.Ratio", by="Elite", ax=ax);
```







```
[35]: fig, ax = subplots(figsize=(8, 8))  
College.boxplot("perc.alumni", by="Elite", ax=ax);
```




```
[36]: fig, ax = subplots(figsize=(8, 8))  
College.boxplot("perc.alumni", by="Private", ax=ax);
```




```
[37]: fig, ax = subplots(figsize=(8, 8))  
College.boxplot("Expend", by="Private", ax=ax);
```




```
[38]: fig, ax = subplots(figsize=(8, 8))  
College.boxplot("Expend", by="Elite", ax=ax);
```



```
[39]: fig, ax = subplots(figsize=(8, 8))  
College.boxplot("Grad.Rate", by="Elite", ax=ax);
```



```
[40]: fig, ax = subplots(figsize=(8, 8))  
College.boxplot("Grad.Rate", by="Private", ax=ax);
```

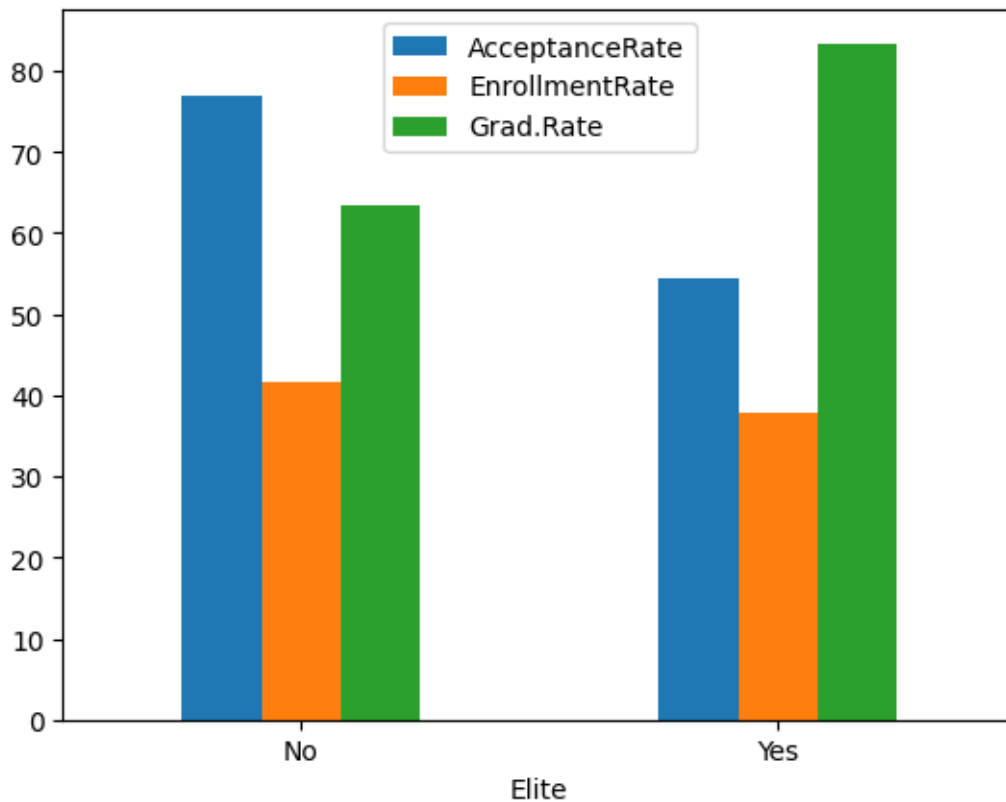


```
[41]: mean_grad_rate = College.groupby("Elite", observed=True)[
      ["AcceptanceRate", "EnrollmentRate", "Grad.Rate"]
    ].mean()
mean_grad_rate
```

```
[41]:
```

	AcceptanceRate	EnrollmentRate	Grad.Rate
Elite			
No	76.963834	41.586452	63.463519
Yes	54.340128	37.748205	83.384615

```
[42]: mean_grad_rate.plot(
      y=["AcceptanceRate", "EnrollmentRate", "Grad.Rate"], kind="bar", rot=0
    );
```



```
[43]: Auto = pd.read_csv("Auto.csv", na_values={"?"})
      print(Auto.shape)
      np.unique(Auto["horsepower"])
```

(397, 9)

```
[43]: array([ 46.,  48.,  49.,  52.,  53.,  54.,  58.,  60.,  61.,  62.,  63.,
        64.,  65.,  66.,  67.,  68.,  69.,  70.,  71.,  72.,  74.,  75.,
        76.,  77.,  78.,  79.,  80.,  81.,  82.,  83.,  84.,  85.,  86.,
        87.,  88.,  89.,  90.,  91.,  92.,  93.,  94.,  95.,  96.,  97.,
        98., 100., 102., 103., 105., 107., 108., 110., 112., 113., 115.,
       116., 120., 122., 125., 129., 130., 132., 133., 135., 137., 138.,
       139., 140., 142., 145., 148., 149., 150., 152., 153., 155., 158.,
       160., 165., 167., 170., 175., 180., 190., 193., 198., 200., 208.,
       210., 215., 220., 225., 230., nan])
```



## 0.0.2 Which predictors are quantitative and which are qualitative?

Rename the misleading column name acceleration to timetoacceleration since it's a tad misleading.

```
[44]: Auto["timetoacceleration"] = Auto["acceleration"]
Auto = Auto.drop("acceleration", axis=1)
```

```
[44]:      mpg  cylinders  displacement  horsepower  weight  year  origin  \
0    18.0         8        307.0        130.0   3504    70      1
1    15.0         8        350.0        165.0   3693    70      1
2    18.0         8        318.0        150.0   3436    70      1
3    16.0         8        304.0        150.0   3433    70      1
4    17.0         8        302.0        140.0   3449    70      1
..    ...         ...         ...         ...     ...    ...    ...
392  27.0         4        140.0         86.0   2790    82      1
393  44.0         4         97.0         52.0   2130    82      2
394  32.0         4        135.0         84.0   2295    82      1
395  28.0         4        120.0         79.0   2625    82      1
396  31.0         4        119.0         82.0   2720    82      1

      name  timetoacceleration
0  chevrolet chevelle malibu        12.0
1      buick skylark 320        11.5
2  plymouth satellite        11.0
3      amc rebel sst        12.0
4      ford torino        10.5
..    ...         ...
392      ford mustang gl        15.6
393      vw pickup        24.6
394  dodge rampage        11.6
395  ford ranger        18.6
396  chevy s-10        19.4

[397 rows x 9 columns]
```

```
[45]: Auto = Auto.dropna()
Auto.shape
```

```
[45]: (392, 9)
```

```
[46]: Auto.describe()
```

```
[46]:      mpg  cylinders  displacement  horsepower  weight  \
count  392.000000  392.000000  392.000000  392.000000  392.000000
mean    23.445918    5.471939   194.411990   104.469388  2977.584184
std     7.805007    1.705783   104.644004    38.491160   849.402560
min     9.000000    3.000000    68.000000    46.000000  1613.000000
25%    17.000000    4.000000   105.000000    75.000000  2225.250000
```

50%	22.750000	4.000000	151.000000	93.500000	2803.500000
75%	29.000000	8.000000	275.750000	126.000000	3614.750000
max	46.600000	8.000000	455.000000	230.000000	5140.000000

	year	origin	timetoacceleration
count	392.000000	392.000000	392.000000
mean	75.979592	1.576531	15.541327
std	3.683737	0.805518	2.758864
min	70.000000	1.000000	8.000000
25%	73.000000	1.000000	13.775000
50%	76.000000	1.000000	15.500000
75%	79.000000	2.000000	17.025000
max	82.000000	3.000000	24.800000

```
[47]: Auto["origin"] = Auto.origin.astype("category")
Auto["year"] = Auto.year.astype("category")
Auto["cylinders"] = Auto.cylinders.astype("category")
print(np.unique(Auto["year"]))
print(np.unique(Auto["cylinders"]))
```

```
[70 71 72 73 74 75 76 77 78 79 80 81 82]
[3 4 5 6 8]
```

```
/tmp/ipykernel_14396/803089839.py:1: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
Auto["origin"] = Auto.origin.astype("category")
/tmp/ipykernel_14396/803089839.py:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
Auto["year"] = Auto.year.astype("category")
/tmp/ipykernel_14396/803089839.py:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
Auto["cylinders"] = Auto.cylinders.astype("category")
```

```
[48]: Auto["origin"] = Auto["origin"].cat.rename_categories(
      {1: "American", 2: "European", 3: "Japanese"}
      )
```

```
np.unique(Auto["origin"])
```

```
/tmp/ipykernel_14396/29212070.py:1: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#returning-a-view-versus-a-copy  
Auto["origin"] = Auto["origin"].cat.rename_categories(  

```

```
[48]: array(['American', 'European', 'Japanese'], dtype=object)
```

```
[49]: Auto.head()
```

```
[49]:      mpg cylinders  displacement  horsepower  weight  year  origin \  
0   18.0          8         307.0         130.0   3504   70  American  
1   15.0          8         350.0         165.0   3693   70  American  
2   18.0          8         318.0         150.0   3436   70  American  
3   16.0          8         304.0         150.0   3433   70  American  
4   17.0          8         302.0         140.0   3449   70  American  
  
      name  timetoacceleration  
0  chevrolet chevelle malibu         12.0  
1      buick skylark 320         11.5  
2    plymouth satellite         11.0  
3      amc rebel sst         12.0  
4      ford torino         10.5
```

```
[50]: Auto = Auto.set_index("name")
```

```
[50]:      mpg cylinders  displacement  horsepower  weight  \  
name  
chevrolet chevelle malibu  18.0          8         307.0         130.0   3504  
buick skylark 320         15.0          8         350.0         165.0   3693  
plymouth satellite        18.0          8         318.0         150.0   3436  
amc rebel sst             16.0          8         304.0         150.0   3433  
ford torino              17.0          8         302.0         140.0   3449  
...                        ...      ...      ...      ...      ...  
ford mustang gl          27.0          4         140.0         86.0   2790  
vw pickup                44.0          4          97.0         52.0   2130  
dodge rampage            32.0          4         135.0         84.0   2295  
ford ranger              28.0          4         120.0         79.0   2625  
chevy s-10               31.0          4         119.0         82.0   2720  
  
      year  origin  timetoacceleration  
name  
chevrolet chevelle malibu   70  American         12.0  
buick skylark 320          70  American         11.5
```

plymouth satellite	70	American	11.0
amc rebel sst	70	American	12.0
ford torino	70	American	10.5
...	...	...	...
ford mustang gl	82	American	15.6
vw pickup	82	European	24.6
dodge rampage	82	American	11.6
ford ranger	82	American	18.6
chevy s-10	82	American	19.4

[392 rows x 8 columns]

```
[51]: Auto
```

```
[51]:
```

	mpg	cylinders	displacement	horsepower	weight	\
name						
chevrolet chevelle malibu	18.0	8	307.0	130.0	3504	
buick skylark 320	15.0	8	350.0	165.0	3693	
plymouth satellite	18.0	8	318.0	150.0	3436	
amc rebel sst	16.0	8	304.0	150.0	3433	
ford torino	17.0	8	302.0	140.0	3449	
...	...	...	...	...	...	
ford mustang gl	27.0	4	140.0	86.0	2790	
vw pickup	44.0	4	97.0	52.0	2130	
dodge rampage	32.0	4	135.0	84.0	2295	
ford ranger	28.0	4	120.0	79.0	2625	
chevy s-10	31.0	4	119.0	82.0	2720	

	year	origin	timetoacceleration
name			
chevrolet chevelle malibu	70	American	12.0
buick skylark 320	70	American	11.5
plymouth satellite	70	American	11.0
amc rebel sst	70	American	12.0
ford torino	70	American	10.5
...	...	...	...
ford mustang gl	82	American	15.6
vw pickup	82	European	24.6
dodge rampage	82	American	11.6
ford ranger	82	American	18.6
chevy s-10	82	American	19.4

[392 rows x 8 columns]

```
[52]: Auto_new = Auto.drop(Auto.index[10:86])
Auto_new.describe()
```

```
[52]:
```

	mpg	displacement	horsepower	weight	timetoacceleration
count	282.000000	282.000000	282.000000	282.000000	282.000000
mean	25.006028	180.120567	99.039007	2884.939716	15.713121
std	7.921384	96.164263	34.197280	793.236373	2.601575
min	11.000000	68.000000	46.000000	1755.000000	8.500000
25%	18.125000	98.000000	74.250000	2188.500000	14.000000
50%	24.500000	140.000000	90.000000	2715.500000	15.500000
75%	31.000000	250.000000	112.000000	3435.250000	17.275000
max	46.600000	455.000000	230.000000	4952.000000	24.600000

```
[53]: Auto_new
```

```
[53]:
```

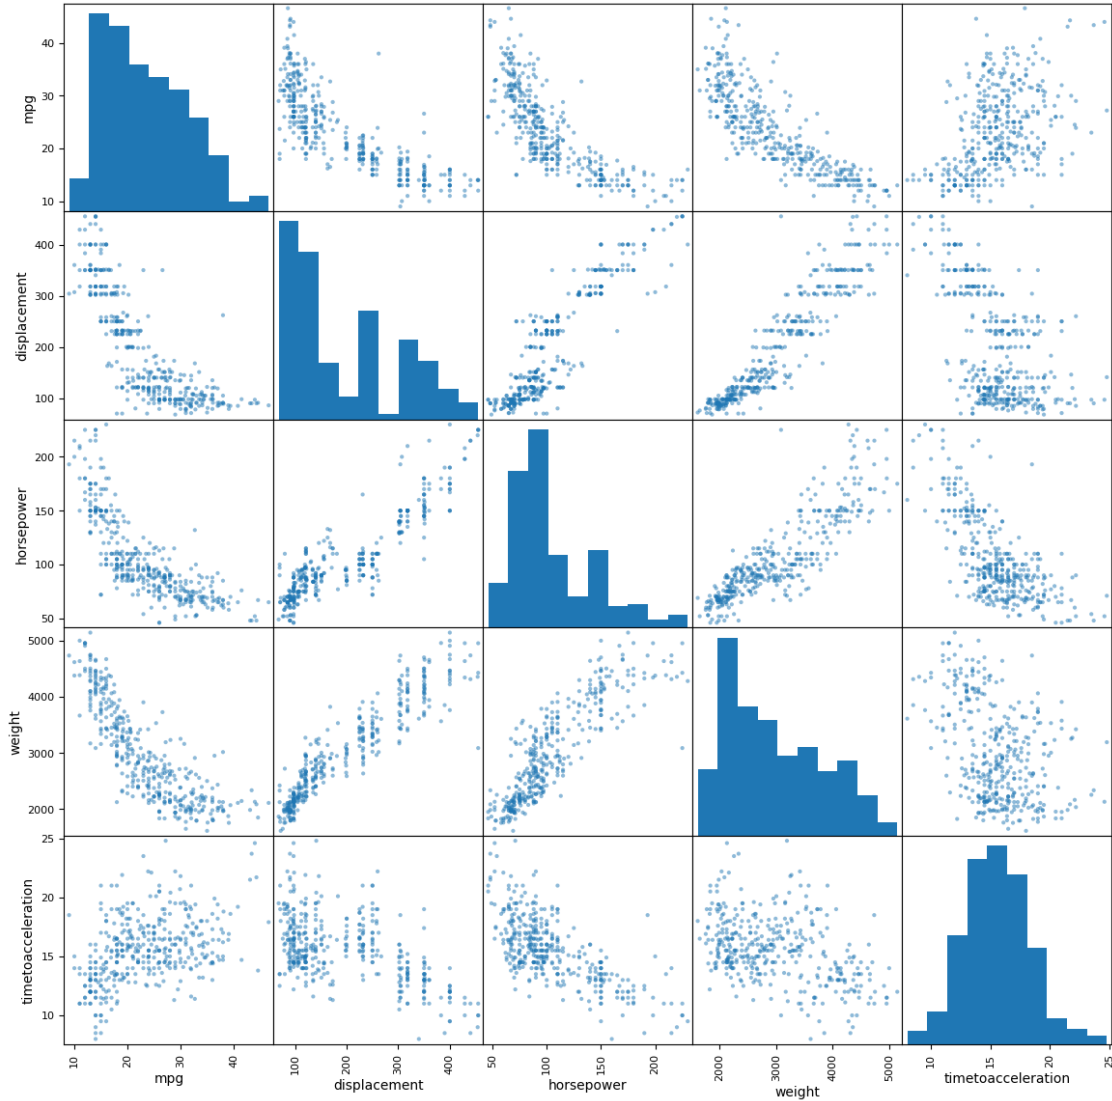
	mpg	cylinders	displacement	horsepower	weight	year	\
name							
buick skylark 320	15.0	8	350.0	165.0	3693	70	
plymouth satellite	18.0	8	318.0	150.0	3436	70	
amc rebel sst	16.0	8	304.0	150.0	3433	70	
ford torino	17.0	8	302.0	140.0	3449	70	
amc ambassador dpl	15.0	8	390.0	190.0	3850	70	
...	...	...	...	...	...	...	
ford mustang gl	27.0	4	140.0	86.0	2790	82	
vw pickup	44.0	4	97.0	52.0	2130	82	
dodge rampage	32.0	4	135.0	84.0	2295	82	
ford ranger	28.0	4	120.0	79.0	2625	82	
chevy s-10	31.0	4	119.0	82.0	2720	82	

	origin	timetoacceleration
name		
buick skylark 320	American	11.5
plymouth satellite	American	11.0
amc rebel sst	American	12.0
ford torino	American	10.5
amc ambassador dpl	American	8.5
...	...	...
ford mustang gl	American	15.6
vw pickup	European	24.6
dodge rampage	American	11.6
ford ranger	American	18.6
chevy s-10	American	19.4

```
[282 rows x 8 columns]
```

Using the full data set, investigate the predictors graphically, using scatter plots or other tools of your choice. Create some plots highlighting the relationships among the predictors. Comment on your findings.

```
[54]: pd.plotting.scatter_matrix(Auto, figsize=(14, 14));
```



### 0.0.3 Findings:

1. Weight and displacement seem to be negatively correlated with MPG.
2. timetoacceleration (0–60 mph in seconds) seems to be positively correlated with MPG. As time to acceleration increases, MPG also increases. The longer the time to acceleration, the better the fuel efficiency.
3. weight is also positively correlated with displacement. As weight increases, so does displacement, i.e., as the body weight increases, so does displacement need to increase.
4. Displacement is seen to increase as the number of cylinders increase. This is expected since displacement is a function of the number of cylinders, amongst other components.

We can conclude that MPG can be predicted using the variables weight, displacement and time-toacceleration.

```
[55]: mean_mpg_origin = Auto.groupby(["origin"], observed=True)[["mpg"]].mean()
mean_mpg_origin
```

```
[55]:          mpg
origin
American  20.033469
European  27.602941
Japanese  30.450633
```

```
[56]: mean_mpg_year = Auto.groupby(["year"], observed=True)[["mpg"]].mean()
mean_mpg_year
```

```
[56]:          mpg
year
70    17.689655
71    21.111111
72    18.714286
73    17.100000
74    22.769231
75    20.266667
76    21.573529
77    23.375000
78    24.061111
79    25.093103
80    33.803704
81    30.185714
82    32.000000
```

```
[57]: mean_mpg_cylinders = Auto.groupby(["cylinders"], observed=True)[["mpg"]].mean()
mean_mpg_cylinders
```

```
[57]:          mpg
cylinders
3        20.550000
4        29.283920
5        27.366667
6        19.973494
8        14.963107
```

We can also observe that fuel efficiency is affected by the make of the car. Japanese > European > American. The year also plays a significant role. Later model cars are more fuel efficient than the earlier models. Cars are also more fuel efficient with lesser number of cylinders. These can also be used as predictors to deduce the MPG.

```
[58]: Boston = load_data("Boston")
Boston.columns
```

```
[58]: Index(['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax',
          'ptratio', 'lstat', 'medv'],
          dtype='object')
```

```
[59]: Boston.shape
```

```
[59]: (506, 13)
```

The rows represent data for 506 suburbs in Boston. The columns represent housing values and variables of interest that may predict housing values in each suburb.

```
[60]: Boston.describe()
```

```
[60]:
```

	crim	zn	indus	chas	nox	rm	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	

	age	dis	rad	tax	ptratio	lstat	\
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	
mean	68.574901	3.795043	9.549407	408.237154	18.455534	12.653063	
std	28.148861	2.105710	8.707259	168.537116	2.164946	7.141062	
min	2.900000	1.129600	1.000000	187.000000	12.600000	1.730000	
25%	45.025000	2.100175	4.000000	279.000000	17.400000	6.950000	
50%	77.500000	3.207450	5.000000	330.000000	19.050000	11.360000	
75%	94.075000	5.188425	24.000000	666.000000	20.200000	16.955000	
max	100.000000	12.126500	24.000000	711.000000	22.000000	37.970000	

	medv
count	506.000000
mean	22.532806
std	9.197104
min	5.000000
25%	17.025000
50%	21.200000
75%	25.000000
max	50.000000

```
[61]: Boston_quant = Boston.drop("chas", axis=1)
```

```
[61]:
```

	crim	zn	indus	nox	rm	age	dis	rad	tax	ptratio	\
0	0.00632	18.0	2.31	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0.469	6.421	78.9	4.9671	2	242	17.8	



2	0.02729	0.0	7.07	0.469	7.185	61.1	4.9671	2	242	17.8
3	0.03237	0.0	2.18	0.458	6.998	45.8	6.0622	3	222	18.7
4	0.06905	0.0	2.18	0.458	7.147	54.2	6.0622	3	222	18.7
..	...	...	...	...	...	...	...	...	...	...
501	0.06263	0.0	11.93	0.573	6.593	69.1	2.4786	1	273	21.0
502	0.04527	0.0	11.93	0.573	6.120	76.7	2.2875	1	273	21.0
503	0.06076	0.0	11.93	0.573	6.976	91.0	2.1675	1	273	21.0
504	0.10959	0.0	11.93	0.573	6.794	89.3	2.3889	1	273	21.0
505	0.04741	0.0	11.93	0.573	6.030	80.8	2.5050	1	273	21.0

	lstat	medv
0	4.98	24.0
1	9.14	21.6
2	4.03	34.7
3	2.94	33.4
4	5.33	36.2
..	...	...
501	9.67	22.4
502	9.08	20.6
503	5.64	23.9
504	6.48	22.0
505	7.88	11.9

[506 rows x 12 columns]

```
[62]: print(np.unique(Boston_quant["zn"]))
median_medv = Boston_quant.groupby(["zn"], observed=True)[["medv"]].median()
median_medv
```

```
[ 0.  12.5  17.5  18.  20.  21.  22.  25.  28.  30.  33.  34.
 35.  40.  45.  52.5 55.  60.  70.  75.  80.  82.5 85.  90.
 95. 100.]
```

```
[62]:      medv
zn
0.0    19.75
12.5   19.90
17.5   33.00
18.0   24.00
20.0   35.20
21.0   21.95
22.0   24.45
25.0   23.10
28.0   22.90
30.0   22.75
33.0   30.90
34.0   26.40
35.0   19.40
```

```


40.0    32.00
45.0    33.45
52.5    23.20
55.0    23.90
60.0    26.60
70.0    24.80
75.0    30.80
80.0    24.50
82.5    33.20
85.0    23.90
90.0    35.40
95.0    41.70
100.0   31.60

```

```

[63]: sns.catplot(
      data=Boston_quant,
      x="zn",
      y="medv",
      kind="bar",
      height=10,
      aspect=2,
      estimator=median,
    );

```



```

[64]: print(np.unique(Boston_quant["rad"]))
      mean_rad = Boston_quant.groupby(["rad"], observed=True)[["medv"]].mean()
      mean_rad

```


```

[ 1  2  3  4  5  6  7  8 24]

```

```
[64]:          medv
rad
1      24.365000
2      26.833333
3      27.928947
4      21.387273
5      25.706957
6      20.976923
7      27.105882
8      30.358333
24     16.403788
```


```
[65]: sns.catplot(data=Boston_quant, x="rad", y="medv", kind="bar", height=10,
↪ aspect=2);
```



```
[66]: sns.set_theme(style="ticks")
g = sns.pairplot(Boston_quant, height=5, aspect=2, diag_kind="kde",
↪ y_vars=["medv"]);
```

Executing <Handle BaseSelectorEventLoop.\_read\_from\_self() created at /usr/lib/python3.12/asyncio/selector\_events.py:280> took 0.289 seconds  
 Executing <Handle IOLoop.\_run\_callback(func tools.par...7dd251c1b880>)) created at /home/linus/ISLP/islpenv/lib/python3.12/site-packages/tornado/platform/asyncio.py:235> took 0.111 seconds  
 Executing <Task pending name='Task-2' coro=<Kernel.poll\_control\_queue() running at /home/linus/ISLP/islpenv/lib/python3.12/site-packages/ipykernel/kernelbase.py:304> wait\_for=<Future pending cb=[Task.task\_wakeup()] created at /home/linus/ISLP/islpenv/lib/python3.12/site-packages/tornado/queues.py:248> cb=[\_chain\_future.<locals>.\_call\_set\_state() at

```
/usr/lib/python3.12/asyncio/futures.py:394] created at  
/usr/lib/python3.12/asyncio/tasks.py:695> took 0.159 seconds
```



Plotting the other quantitative columns against medv (Median value of owner-occupied homes), we can see that: 1. crim is negatively correlated with medv. i.e., as crime rate increases, median value of homes decrease. 2. indus is negatively correlated with medv which is expected as industrialisation of a town increases, the house prices decrease. 3. nox is negatively correlated with medv which is also expected. 4. as the number of rooms (rm) increase, so does the value of the home. 5. as the proportion of homes built prior to 1940 increase, the value of homes in that area decrease. There are some notable outliers, but that appears to be the general trend. 6. There is a clear relationship in the lstat (lower status of population percent) versus medv where medv decreases with the increase in lstat on the x-axis.


```
[67]: Boston_quant["zn"].value_counts()
```

```
[67]: zn
0.0      372
20.0      21
80.0      15
22.0      10
12.5      10
25.0      10
40.0       7
45.0       6
30.0       6
90.0       5
95.0       4
60.0       4
21.0       4
33.0       4
55.0       3
70.0       3
34.0       3
52.5       3
35.0       3
28.0       3
75.0       3
82.5       2
85.0       2
17.5       1
100.0      1
18.0       1
Name: count, dtype: int64
```


```
[68]: g = sns.pairplot(Boston_quant, height=5, aspect=2, diag_kind="kde",
↳ y_vars=["crim"]);
```



```
[69]: sns.catplot(data=Boston_quant, x="zn", y="crim", kind="bar", height=6,
↳ aspect=2);
```




```
[70]: sns.displot(data=Boston_quant, x="indus", y="crim", height=4, aspect=3);
```




```
[71]: sns.displot(data=Boston_quant, x="age", y="crim", height=4, aspect=3);
```

Executing <Handle BaseAsyncIOLoop.\_handle\_events(28, 1) created at


```
/usr/lib/python3.12/asyncio/selector_events.py:280> took 0.290 seconds  
Executing <Handle IOLoop._run_callback(functools.par...7dd2511184a0>)) created  
at /home/linus/ISLP/islpenv/lib/python3.12/site-  
packages/tornado/platform/asyncio.py:235> took 0.299 seconds
```




```
[72]: sns.displot(data=Boston_quant, x="nox", y="crim", height=4, aspect=3);
```




```
[73]: sns.displot(data=Boston_quant, x="dis", y="crim", height=4, aspect=3);
```



```
[74]: sns.displot(data=Boston_quant, x="tax", y="crim", height=4, aspect=3);
```



```
[75]: sns.displot(data=Boston_quant, x="ptratio", y="crim", height=4, aspect=3);
```



We've already seen that there appears to be a relationship b/w crime rate and medv where a higher crime rate is associated with lower property prices. Additionally, plotting the other quantitative variables against crime rate (crim), we can perceive the following:

1. No-zoned areas or towns are associated with higher crime rate compared to all other zoning percentages.
2. For some reason, industrialization of around 18% displays a spike in the crime rate compared to the other suburbs. This might be worth investigating further.
3. Suburbs with  $\text{nox} > 0.55$  or so have an elevated crime rate. That could be because lower strata income people live in those areas, and they are more inclined to criminal activities.
4. There also seems to be an increasing relationship b/w crime rate and percentage of homes built prior to 1940. Once that percentage crosses 40%, there is an increasing number of suburbs that exhibit elevated crime rates.
5. Suburbs within a distance to Boston employment centres that range from 1 to 4.5 show an elevated crime rate. This needs to be investigated further. Where are these employment centres located?
6. There seems to be a higher incidence of crimes for areas with tax rate around 670. Why?
7. The crime rate does not seem to have a strong relationship with ptratio, but for around point 20.1 where the crime rate spikes compared to the other areas.
8. Crime rate decreases as the median value of properties rise across

suburbs as a whole.

*Do any of the suburbs of Boston appear to have particularly high crime rates? Tax rates? Pupil-teacher ratios? Comment on the range of each predictor.*

```
[76]: Boston_crim = Boston_quant.sort_values(
        by="crim", axis=0, ascending=False, inplace=False
    )
    top_crim = Boston_crim.head()
    print(top_crim)
    df = pd.DataFrame((Boston_quant.min(), Boston_quant.max()), index=["Min",
        ↪ "Max"])
    df
```

	crim	zn	indus	nox	rm	age	dis	rad	tax	ptratio	\
380	88.9762	0.0	18.1	0.671	6.968	91.9	1.4165	24	666	20.2	
418	73.5341	0.0	18.1	0.679	5.957	100.0	1.8026	24	666	20.2	
405	67.9208	0.0	18.1	0.693	5.683	100.0	1.4254	24	666	20.2	
410	51.1358	0.0	18.1	0.597	5.757	100.0	1.4130	24	666	20.2	
414	45.7461	0.0	18.1	0.693	4.519	100.0	1.6582	24	666	20.2	

	lstat	medv
380	17.21	10.4
418	20.62	8.8
405	22.98	5.0
410	10.11	15.0
414	36.98	7.0

```
[76]:
```

	crim	zn	indus	nox	rm	age	dis	rad	tax	\
Min	0.00632	0.0	0.46	0.385	3.561	2.9	1.1296	1.0	187.0	
Max	88.97620	100.0	27.74	0.871	8.780	100.0	12.1265	24.0	711.0	

	ptratio	lstat	medv
Min	12.6	1.73	5.0
Max	22.0	37.97	50.0

As we can see from the dataset above, the top five crime rate suburbs are unzoned, have an industrialization rate of 18.1%, nox of 0.671, rooms ranging from 4.5 to 7, percentage of houses built prior to 1940 ranging from 92 to 100%, high tax rate of 666per10,000 property tax , ptratio of 20.2. The lstat varies from 10.11 rto 36.98 and the median house values from 5.0 to 15.0 which are among the lowest. The index of accessibility to radial highways is 24 which is the best rank amongst all the suburbs.

*How many of the suburbs in this data set bound the Charles River?*

```
[77]: len(Boston.query("chas == 1"))
```

```
[77]: 35
```

*What is the median pupil-teacher ratio among the towns in this data set?*



```
[78]: median(Boston_quant["ptratio"])
```

```
[78]: 19.05
```

*Which suburb of Boston has lowest median value of owner-occupied homes? What are the values of the other predictors for that suburb, and how do those values compare to the overall ranges for those predictors? Comment on your findings.*

```
[79]: lowest_medv = Boston_quant[Boston_quant["medv"] == Boston_quant["medv"].min()]
lowest_medv
```

```
[79]:      crim    zn  indus    nox    rm    age    dis  rad  tax  ptratio  \
398  38.3518  0.0   18.1  0.693  5.453  100.0  1.4896   24  666    20.2
405  67.9208  0.0   18.1  0.693  5.683  100.0  1.4254   24  666    20.2

      lstat  medv
398   30.59    5.0
405   22.98    5.0
```

From the above two data points with the lowest median value for owner-occupied homes, it's evident that crime rate by itself does not determine the median value of homes for those regions. Except for lstat and crim, the other predictors match exactly the two data points. lstat for these two data points are high at 22.98 and 30.59 respectively. In these two cases, the other predictors do a better job of explaining the median value for the homes in these suburbs or these suburbs are neighbouring each other.

*In this data set, how many of the suburbs average more than seven rooms per dwelling? More than eight rooms per dwelling? Comment on the suburbs that average more than eight rooms per dwelling.*

```
[80]: len(Boston_quant[Boston_quant["rm"] > 7])
```

```
[80]: 64
```

```
[81]: eight_rooms = Boston_quant[Boston_quant["rm"] > 8]
print(len(eight_rooms))
eight_rooms
```

```
13
```

```
[81]:      crim    zn  indus    nox    rm    age    dis  rad  tax  ptratio  \
97   0.12083  0.0   2.89  0.4450  8.069  76.0  3.4952   2  276    18.0
163  1.51902  0.0  19.58  0.6050  8.375  93.9  2.1620   5  403    14.7
204  0.02009 95.0   2.68  0.4161  8.034  31.9  5.1180   4  224    14.7
224  0.31533  0.0   6.20  0.5040  8.266  78.3  2.8944   8  307    17.4
225  0.52693  0.0   6.20  0.5040  8.725  83.0  2.8944   8  307    17.4
226  0.38214  0.0   6.20  0.5040  8.040  86.5  3.2157   8  307    17.4
232  0.57529  0.0   6.20  0.5070  8.337  73.3  3.8384   8  307    17.4
233  0.33147  0.0   6.20  0.5070  8.247  70.4  3.6519   8  307    17.4
253  0.36894 22.0   5.86  0.4310  8.259   8.4  8.9067   7  330    19.1
```

257	0.61154	20.0	3.97	0.6470	8.704	86.9	1.8010	5	264	13.0
262	0.52014	20.0	3.97	0.6470	8.398	91.5	2.2885	5	264	13.0
267	0.57834	20.0	3.97	0.5750	8.297	67.0	2.4216	5	264	13.0
364	3.47428	0.0	18.10	0.7180	8.780	82.9	1.9047	24	666	20.2

	lstat	medv
97	4.21	38.7
163	3.32	50.0
204	2.88	50.0
224	4.14	44.8
225	4.63	50.0
226	3.13	37.6
232	2.47	41.7
233	3.95	48.3
253	3.54	42.8
257	5.12	50.0
262	5.91	48.8
267	7.44	50.0
364	5.29	21.9

There are 13 suburbs that average more than 8 rooms per dwelling.

The median value for these dwellings range from 21.9 to 50.0 which is the priciest.

The crime rate in these suburbs is extremely low with the highest at around 3.5%.

Industrialization of these suburbs is also low with 19.58 the maximum.

The percentage of people from the lower income strata tops out at 7.44%

A substantial percentage of dwellings are built prior to 1940 which could explain the higher number of rooms with only one outlier at 8.4%.

```
[82]: allDone()
```

```
<IPython.lib.display.Audio object>
```