

Labs

February 21, 2025

1 Lab: Linear Regression

1.1 Set up IPython libraries for customizing notebook display

```
[1]: from notebookfuncs import *
```

1.2 Import standard libraries

```
[2]: import numpy as np
import pandas as pd
from matplotlib.pyplot import subplots
```

1.3 New imports

```
[3]: import statsmodels.api as sm
```

1.4 Import statsmodels objects

```
[4]: from statsmodels.stats.outliers_influence import variance_inflation_factor as VIF
↪ VIF
from statsmodels.stats.anova import anova_lm
```

1.5 Import ISLP objects

```
[5]: from ISLP import load_data
from ISLP.models import ModelSpec as MS, summarize, poly
```

1.6 Import User Functions

```
[6]: from userfuncs import *
```

1.7 Inspecting objects and namespaces

```
[7]: dir()
```

```
[7]: ['Audio',
      'In',
```

```
'InteractiveShell',
'Latex',
'MS',
'Markdown',
'Math',
'Out',
'VIF',
'_',
'__',
'___',
'__builtin__',
'__builtins__',
'__doc__',
'__loader__',
'__name__',
'__package__',
'__spec__',
'_dh',
'_i',
'_i1',
'_i2',
'_i3',
'_i4',
'_i5',
'_i6',
'_i7',
'_ih',
'_ii',
'_iii',
'_oh',
'allDone',
'anova_lm',
'calculate_VIFs',
'check_symmetric',
'display',
'display_DFFITS_plot',
'display_cooks_distance_plot',
'display_hat_leverage_cutoffs',
'display_hat_leverage_plot',
'display_residuals_plot',
'display_studentized_residuals',
'dmatrices',
'exit',
'get_influence_points',
'get_ipython',
'get_results_df',
'identify_highest_VIF_feature',
```

```

'identify_least_significant_feature',
'influence_plot',
'is_numeric_dtype',
'is_pos_def',
'is_symmetric_pos_def',
'load_data',
'np',
'open',
'pd',
'perform_analysis',
'poly',
'printlatex',
'printmd',
'px',
'quit',
'sm',
'smf',
'standardize',
'stats',
'subplots',
'summarize']

```

```

[8]: A = np.array([3, 5, 11])
     dir(A)

```

```

[8]: ['T',
      '__abs__',
      '__add__',
      '__and__',
      '__array__',
      '__array_finalize__',
      '__array_function__',
      '__array_interface__',
      '__array_prepare__',
      '__array_priority__',
      '__array_struct__',
      '__array_ufunc__',
      '__array_wrap__',
      '__bool__',
      '__buffer__',
      '__class__',
      '__class_getitem__',
      '__complex__',
      '__contains__',
      '__copy__',
      '__deepcopy__',
      '__delattr__',

```

```
'__delitem__',
'__dir__',
'__divmod__',
'__dlpack__',
'__dlpack_device__',
'__doc__',
'__eq__',
'__float__',
'__floordiv__',
'__format__',
'__ge__',
'__getattr__',
'__getitem__',
'__getstate__',
'__gt__',
'__hash__',
'__iadd__',
'__iand__',
'__ifloordiv__',
'__ilshift__',
'__imatmul__',
'__imod__',
'__imul__',
'__index__',
'__init__',
'__init_subclass__',
'__int__',
'__invert__',
'__ior__',
'__ipow__',
'__irshift__',
'__isub__',
'__iter__',
'__itruediv__',
'__ixor__',
'__le__',
'__len__',
'__lshift__',
'__lt__',
'__matmul__',
'__mod__',
'__mul__',
'__ne__',
'__neg__',
'__new__',
'__or__',
'__pos__',
```

```
'__pow__',
'__radd__',
'__rand__',
'__rdivmod__',
'__reduce__',
'__reduce_ex__',
'__repr__',
'__rfloordiv__',
'__rlshift__',
'__rmatmul__',
'__rmod__',
'__rmul__',
'__ror__',
'__rpow__',
'__rrshift__',
'__rshift__',
'__rsub__',
'__rtruediv__',
'__rxor__',
'__setattr__',
'__setitem__',
'__setstate__',
'__sizeof__',
'__str__',
'__sub__',
'__subclasshook__',
'__truediv__',
'__xor__',
'all',
'any',
'argmax',
'argmin',
'argpartition',
'argsort',
'astype',
'base',
'byteswap',
'choose',
'clip',
'compress',
'conj',
'conjugate',
'copy',
'ctypes',
'cumprod',
'cumsum',
'data',
```

'diagonal',
'dot',
'dtype',
'dump',
'dumps',
'fill',
'flags',
'flat',
'flatten',
'getfield',
'imag',
'item',
'itemset',
'itemsizes',
'max',
'mean',
'min',
'nbytes',
'ndim',
'newbyteorder',
'nonzero',
'partition',
'prod',
'ptp',
'put',
'ravel',
'real',
'repeat',
'reshape',
'resize',
'round',
'searchsorted',
'setfield',
'setflags',
'shape',
'size',
'sort',
'squeeze',
'std',
'strides',
'sum',
'swapaxes',
'take',
'tobytes',
'tofile',
'tolist',
'tostring',

```
'trace',  
'transpose',  
'var',  
'view']
```

```
[9]: A.sum()
```

```
[9]: 19
```

1.8 Simple Linear Regression

1.8.1 We will use the Boston housing dataset which is in the package ISLP

```
[10]: Boston = load_data("Boston")  
Boston.columns
```

```
[10]: Index(['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax',  
        'ptratio', 'lstat', 'medv'],  
        dtype='object')
```

```
[11]: len(Boston.columns)
```

```
[11]: 13
```

```
[12]: Boston?
```

1.8.2 Use sm.OLS to fit a simple linear regression

```
[13]: X = pd.DataFrame({"intercept": np.ones(Boston.shape[0]), "lstat":  
        ↪Boston["lstat"]})  
X.head()
```

```
[13]:   intercept  lstat  
0         1.0    4.98  
1         1.0    9.14  
2         1.0    4.03  
3         1.0    2.94  
4         1.0    5.33
```

1.8.3 Extract the response and fit the model.

```
[14]: y = Boston["medv"]  
model = sm.OLS(y, X)  
results = model.fit()
```

```
[14]: <statsmodels.regression.linear_model.RegressionResultsWrapper at 0x7370ae75f200>
```

1.8.4 Summarize the results using the ISLP method summarize

```
[15]: summarize(results)
```

```
[15]:      coef  std err      t  P>|t|
intercept  34.5538    0.563  61.415   0.0
lstat      -0.9500    0.039 -24.528   0.0
```

1.9 Using Transformations: Fit and Transform

```
[16]: design = MS(["lstat"])
design = design.fit(Boston)
X = design.transform(Boston)
X.head()
```

```
[16]:      intercept  lstat
0          1.0    4.98
1          1.0    9.14
2          1.0    4.03
3          1.0    2.94
4          1.0    5.33
```

```
[17]: design = MS(["lstat"])
design = design.fit_transform(Boston)
X.head()
```

```
[17]:      intercept  lstat
0          1.0    4.98
1          1.0    9.14
2          1.0    4.03
3          1.0    2.94
4          1.0    5.33
```

1.9.1 Full and exhaustive summary of the fit

```
[18]: results.summary()
```

```
[18]:
```

Dep. Variable:	medv	R-squared:	0.544
Model:	OLS	Adj. R-squared:	0.543
Method:	Least Squares	F-statistic:	601.6
Date:	Fri, 21 Feb 2025	Prob (F-statistic):	5.08e-88
Time:	19:21:12	Log-Likelihood:	-1641.5
No. Observations:	506	AIC:	3287.
Df Residuals:	504	BIC:	3295.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
intercept	34.5538	0.563	61.415	0.000	33.448	35.659
lstat	-0.9500	0.039	-24.528	0.000	-1.026	-0.874
Omnibus:		137.043	Durbin-Watson:		0.892	
Prob(Omnibus):		0.000	Jarque-Bera (JB):		291.373	
Skew:		1.453	Prob(JB):		5.36e-64	
Kurtosis:		5.319	Cond. No.		29.7	

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

1.9.2 Fitted coefficients can be retrieved as the *params* attribute of results

```
[19]: results.params
```

```
[19]: intercept    34.553841
      lstat        -0.950049
      dtype: float64
```

1.9.3 Computing predictions

```
[20]: design = MS(["lstat"])
      new_df = pd.DataFrame({"lstat": [5, 10, 15]})
      print(new_df)
      design = design.fit(new_df)
      newX = design.transform(new_df)
      newX
```

```
      lstat
0         5
1        10
2        15
```

```
[20]:      intercept  lstat
0         1.0        5
1         1.0       10
2         1.0       15
```

```
[21]: new_predictions = results.get_prediction(newX)
      new_predictions.predicted_mean
```

```
[21]: array([29.80359411, 25.05334734, 20.30310057])
```

1.9.4 We can predict confidence intervals for the predicted values.

```
[22]: new_predictions.conf_int(alpha=0.05)
```

```
[22]: array([[29.00741194, 30.59977628],
           [24.47413202, 25.63256267],
           [19.73158815, 20.87461299]])
```

1.9.5 We can obtain prediction intervals for the values which are wider than the confidence intervals since they're for a specific instance of `lstat` by setting `obs=True`.


```
[23]: new_predictions.conf_int(obs=True, alpha=0.05)
```

```
[23]: array([[17.56567478, 42.04151344],
           [12.82762635, 37.27906833],
           [ 8.0777421 , 32.52845905]])
```

1.9.6 Plot `medv` and `lstat` using `DataFrame.plot.scatter()` and add the regression line to the resulting plot.

```
[24]: ax = Boston.plot.scatter("lstat", "medv")
      ax.axline(
          (ax.get_xlim()[0], results.params.iloc[0]),
          slope=results.params.iloc[1],
          color="r",
          linewidth=3,
      )
```

```
[24]: <matplotlib.lines.AxLine at 0x7370a6b55ee0>
```




- There is some evidence of non-linearity in the relationship b/w lstat and medv.

1.9.7 Find the fitted values and residuals of the fit as attributes of the results object as *results.fittedvalues* and *results.resid*.

- The `get_influence()` method computes various influence measures of the regression.

```
[25]: _, ax = subplots(figsize=(8, 8))
      ax.scatter(results.fittedvalues, results.resid)
      ax.set_xlabel("Fitted values")
      ax.set_ylabel("Residuals")
      ax.axhline(0, c="k", ls="--")
```


```
[25]: <matplotlib.lines.Line2D at 0x7370a5733290>
```



- On the basis of the residual plot, there is some evidence of non-linearity.

1.9.8 Leverage statistics can be computed for any number of predictors using the `hat_matrix_diag` attribute of the value returned by the `get_influence()` method.


```
[26]: display_hat_leverage_cutoffs(results)
```




```
[27]: display_hat_leverage_plot(results)
```

```
[28]: display_cooks_distance_plot(results)
```


```
[28]:
```






```
[29]: display_DFFITS_plot(results)
```

```
[29]:
```





```
[30]: inf_df, _ = get_influence_points(results)
      inf_df
```

```
n = 506.0, p = 2
Average Hat Leverage: 0.003952569169960474
Hat Leverage Cutoff = 2 * Average Hat Leverage = 0.007905138339920948
DFBetas Cutoff = 3 / sqrt(n) = 0.1333662673423161
DFFITS Cutoff = 2 * sqrt(p/n) = 0.1257389226923863
Cooks Distance Cutoff = 1.0
Cooks Distance p-value Cutoff = 0.05
Studentized Residuals Cutoff = 3.0
Studentized Residuals p-value Cutoff = 0.01
```

```
[30]:
```

	dfb_intercept	dfb_lstat	cooks_d	hat_diag	student_resid	dffits	\
161	0.226022	-0.189640	0.025315	0.006609	2.776857	0.226503	
162	0.225506	-0.188310	0.025219	0.006450	2.806416	0.226112	
163	0.219862	-0.176384	0.024252	0.005359	3.024654	0.222011	
166	0.217770	-0.172500	0.023921	0.005089	3.084034	0.220566	
186	0.212939	-0.164025	0.023201	0.004589	3.201426	0.217377	

195	0.221577	-0.179716	0.024534	0.005617	2.970018	0.223224
203	0.199749	-0.157618	0.020220	0.005013	2.853126	0.202515
204	0.221985	-0.180535	0.024602	0.005685	2.955977	0.223517
214	-0.197509	0.297576	0.051465	0.013063	2.807647	0.323011
225	0.211641	-0.161830	0.023017	0.004476	3.229640	0.216554
228	0.178647	-0.140414	0.016255	0.004938	2.573814	0.181309
233	0.196791	-0.154507	0.019680	0.004918	2.841931	0.199782
257	0.207834	-0.155541	0.022502	0.004180	3.306527	0.214221
261	0.128039	-0.084194	0.009646	0.003106	2.501385	0.139616
262	0.189166	-0.136023	0.019246	0.003742	3.231093	0.198020
267	0.184377	-0.119477	0.020008	0.003032	3.672331	0.202505
280	0.164168	-0.129770	0.013717	0.005047	2.335776	0.166365
283	0.220671	-0.177936	0.024384	0.005476	2.999671	0.222580
368	0.220170	-0.176971	0.024302	0.005402	3.015284	0.222227
369	0.217594	-0.172182	0.023894	0.005068	3.088724	0.220447
370	0.221623	-0.179808	0.024542	0.005625	2.968458	0.223257
371	0.155509	-0.078029	0.018381	0.002355	4.004703	0.194572
372	0.165278	-0.091836	0.018763	0.002529	3.901020	0.196431
374	-0.294291	0.401657	0.086162	0.026865	2.511537	0.417300
412	-0.253809	0.357605	0.070029	0.020290	2.615542	0.376405

	student_resid_pvalue	hat_influence	cooks_d_pvalue
161	0.002847	0.018353	0.975004
162	0.002602	0.018100	0.975097
163	0.001308	0.016208	0.976041
166	0.001077	0.015694	0.976364
186	0.000727	0.014692	0.977067
195	0.001560	0.016683	0.975766
203	0.002254	0.014302	0.979984
204	0.001632	0.016805	0.975699
214	0.002592	0.036676	0.949842
225	0.000660	0.014455	0.977247
228	0.005172	0.012709	0.983877
233	0.002333	0.013975	0.980513
257	0.000506	0.013821	0.977750
261	0.006344	0.007769	0.990401
262	0.000657	0.012090	0.980939
267	0.000133	0.011133	0.980191
280	0.009947	0.011789	0.986377
283	0.001418	0.016425	0.975912
368	0.001349	0.016290	0.975992
369	0.001061	0.015654	0.976391
370	0.001568	0.016697	0.975758
371	0.000036	0.009431	0.981788
372	0.000054	0.009866	0.981412
374	0.006166	0.067473	0.917459
412	0.004588	0.053070	0.932376

1.9.9 For a more conservative cutoff values for `hat_diag`, we have the following influence point(s):

```
[31]: inf_df[inf_df["hat_diag"] > (3 * np.mean(inf_df["hat_diag"]))]
```

```
[31]:      dfb_intercept  dfb_lstat   cooks_d  hat_diag  student_resid  dffits  \
374      -0.294291    0.401657  0.086162  0.026865      2.511537  0.417300
412      -0.253809    0.357605  0.070029  0.020290      2.615542  0.376405

      student_resid_pvalue  hat_influence  cooks_d_pvalue
374              0.006166      0.067473      0.917459
412              0.004588      0.053070      0.932376
```

1.9.10 Using DFFITS cutoff, we have the following influential points

```
[32]: inf_df[inf_df["dffits"] > 2 * np.sqrt(len(results.params) / results.nobs)]
```

```
[32]:      dfb_intercept  dfb_lstat   cooks_d  hat_diag  student_resid  dffits  \
161      0.226022   -0.189640  0.025315  0.006609      2.776857  0.226503
162      0.225506   -0.188310  0.025219  0.006450      2.806416  0.226112
163      0.219862   -0.176384  0.024252  0.005359      3.024654  0.222011
166      0.217770   -0.172500  0.023921  0.005089      3.084034  0.220566
186      0.212939   -0.164025  0.023201  0.004589      3.201426  0.217377
195      0.221577   -0.179716  0.024534  0.005617      2.970018  0.223224
203      0.199749   -0.157618  0.020220  0.005013      2.853126  0.202515
204      0.221985   -0.180535  0.024602  0.005685      2.955977  0.223517
214     -0.197509    0.297576  0.051465  0.013063      2.807647  0.323011
225      0.211641   -0.161830  0.023017  0.004476      3.229640  0.216554
228      0.178647   -0.140414  0.016255  0.004938      2.573814  0.181309
233      0.196791   -0.154507  0.019680  0.004918      2.841931  0.199782
257      0.207834   -0.155541  0.022502  0.004180      3.306527  0.214221
261      0.128039   -0.084194  0.009646  0.003106      2.501385  0.139616
262      0.189166   -0.136023  0.019246  0.003742      3.231093  0.198020
267      0.184377   -0.119477  0.020008  0.003032      3.672331  0.202505
280      0.164168   -0.129770  0.013717  0.005047      2.335776  0.166365
283      0.220671   -0.177936  0.024384  0.005476      2.999671  0.222580
368      0.220170   -0.176971  0.024302  0.005402      3.015284  0.222227
369      0.217594   -0.172182  0.023894  0.005068      3.088724  0.220447
370      0.221623   -0.179808  0.024542  0.005625      2.968458  0.223257
371      0.155509   -0.078029  0.018381  0.002355      4.004703  0.194572
372      0.165278   -0.091836  0.018763  0.002529      3.901020  0.196431
374     -0.294291    0.401657  0.086162  0.026865      2.511537  0.417300
412     -0.253809    0.357605  0.070029  0.020290      2.615542  0.376405

      student_resid_pvalue  hat_influence  cooks_d_pvalue
161              0.002847      0.018353      0.975004
162              0.002602      0.018100      0.975097
```

163	0.001308	0.016208	0.976041
166	0.001077	0.015694	0.976364
186	0.000727	0.014692	0.977067
195	0.001560	0.016683	0.975766
203	0.002254	0.014302	0.979984
204	0.001632	0.016805	0.975699
214	0.002592	0.036676	0.949842
225	0.000660	0.014455	0.977247
228	0.005172	0.012709	0.983877
233	0.002333	0.013975	0.980513
257	0.000506	0.013821	0.977750
261	0.006344	0.007769	0.990401
262	0.000657	0.012090	0.980939
267	0.000133	0.011133	0.980191
280	0.009947	0.011789	0.986377
283	0.001418	0.016425	0.975912
368	0.001349	0.016290	0.975992
369	0.001061	0.015654	0.976391
370	0.001568	0.016697	0.975758
371	0.000036	0.009431	0.981788
372	0.000054	0.009866	0.981412
374	0.006166	0.067473	0.917459
412	0.004588	0.053070	0.932376

1.9.11 Using Cooks Distance, we have the following influential points

```
[33]: inf_df[inf_df["cooks_d"] > 1.0]
```

```
[33]: Empty DataFrame
Columns: [dfb_intercept, dfb_lstat, cooks_d, hat_diag, student_resid, dffits,
student_resid_pvalue, hat_influence, cooks_d_pvalue]
Index: []
```

1.9.12 Using Cooks Distance p-values, we have the following influential points

```
[34]: inf_df[inf_df["cooks_d_pvalue"] < 0.05]
```

```
[34]: Empty DataFrame
Columns: [dfb_intercept, dfb_lstat, cooks_d, hat_diag, student_resid, dffits,
student_resid_pvalue, hat_influence, cooks_d_pvalue]
Index: []
```

1.9.13 Using DFBeta for intercept, we have the following influential points

```
[35]: inf_df[inf_df["dfb_intercept"] > (3 / np.sqrt(results.nobs))]
```

```

[35]:      dfb_intercept  dfb_lstat   cooks_d  hat_diag  student_resid  dffits  \
161      0.226022 -0.189640  0.025315  0.006609      2.776857  0.226503
162      0.225506 -0.188310  0.025219  0.006450      2.806416  0.226112
163      0.219862 -0.176384  0.024252  0.005359      3.024654  0.222011
166      0.217770 -0.172500  0.023921  0.005089      3.084034  0.220566
186      0.212939 -0.164025  0.023201  0.004589      3.201426  0.217377
195      0.221577 -0.179716  0.024534  0.005617      2.970018  0.223224
203      0.199749 -0.157618  0.020220  0.005013      2.853126  0.202515
204      0.221985 -0.180535  0.024602  0.005685      2.955977  0.223517
225      0.211641 -0.161830  0.023017  0.004476      3.229640  0.216554
228      0.178647 -0.140414  0.016255  0.004938      2.573814  0.181309
233      0.196791 -0.154507  0.019680  0.004918      2.841931  0.199782
257      0.207834 -0.155541  0.022502  0.004180      3.306527  0.214221
262      0.189166 -0.136023  0.019246  0.003742      3.231093  0.198020
267      0.184377 -0.119477  0.020008  0.003032      3.672331  0.202505
280      0.164168 -0.129770  0.013717  0.005047      2.335776  0.166365
283      0.220671 -0.177936  0.024384  0.005476      2.999671  0.222580
368      0.220170 -0.176971  0.024302  0.005402      3.015284  0.222227
369      0.217594 -0.172182  0.023894  0.005068      3.088724  0.220447
370      0.221623 -0.179808  0.024542  0.005625      2.968458  0.223257
371      0.155509 -0.078029  0.018381  0.002355      4.004703  0.194572
372      0.165278 -0.091836  0.018763  0.002529      3.901020  0.196431

      student_resid_pvalue  hat_influence  cooks_d_pvalue
161      0.002847      0.018353      0.975004
162      0.002602      0.018100      0.975097
163      0.001308      0.016208      0.976041
166      0.001077      0.015694      0.976364
186      0.000727      0.014692      0.977067
195      0.001560      0.016683      0.975766
203      0.002254      0.014302      0.979984
204      0.001632      0.016805      0.975699
225      0.000660      0.014455      0.977247
228      0.005172      0.012709      0.983877
233      0.002333      0.013975      0.980513
257      0.000506      0.013821      0.977750
262      0.000657      0.012090      0.980939
267      0.000133      0.011133      0.980191
280      0.009947      0.011789      0.986377
283      0.001418      0.016425      0.975912
368      0.001349      0.016290      0.975992
369      0.001061      0.015654      0.976391
370      0.001568      0.016697      0.975758
371      0.000036      0.009431      0.981788
372      0.000054      0.009866      0.981412

```

1.9.14 Using DFBeta for lstat, we have the following influential points

```
[36]: inf_df[inf_df["dfb_lstat"] > (3 / np.sqrt(results.nobs))]
```


```
[36]:      dfb_intercept  dfb_lstat  cooks_d  hat_diag  student_resid  dffits  \
214      -0.197509   0.297576  0.051465  0.013063      2.807647  0.323011
374      -0.294291   0.401657  0.086162  0.026865      2.511537  0.417300
412      -0.253809   0.357605  0.070029  0.020290      2.615542  0.376405

      student_resid_pvalue  hat_influence  cooks_d_pvalue
214              0.002592        0.036676        0.949842
374              0.006166        0.067473        0.917459
412              0.004588        0.053070        0.932376
```

1.9.15 Multiple linear regression


```
[37]: Boston.plot.scatter("age", "medv")
X = MS(["lstat", "age"]).fit_transform(Boston)
model1 = sm.OLS(y, X)
results1 = model1.fit()
summarize(results1)
```

```
[37]:      coef  std err      t  P>|t|
intercept  33.2228    0.731  45.458  0.000
lstat      -1.0321    0.048 -21.416  0.000
age         0.0345    0.012   2.826  0.005
```



```
[38]: Boston["logage"] = np.log(Boston["age"])
Boston.plot.scatter("logage", "medv")
X = MS(["lstat", "logage"]).fit_transform(Boston)
model1 = sm.OLS(y, X)
resultslog = model1.fit()
print(summarize(resultslog))
```


	coef	std err	t	P> t
intercept	30.2143	1.947	15.517	0.00
lstat	-1.0051	0.045	-22.213	0.00
logage	1.2312	0.529	2.327	0.02



```
[39]: Boston["sqrtage"] = np.sqrt(Boston["age"])
Boston.plot.scatter("sqrtage", "medv")
X = MS(["lstat", "sqrtage"]).fit_transform(Boston)
model1 = sm.OLS(y, X)
resultssqrt = model1.fit()
summarize(resultssqrt)
```

```
[39]:
```

	coef	std err	t	P> t
intercept	31.8635	1.174	27.139	0.000
lstat	-1.0203	0.047	-21.703	0.000
sqrtage	0.4450	0.171	2.606	0.009



```
[40]: Boston = Boston.drop(columns=["logage", "sqrtage"])
```

```
[40]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	
..	
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273	
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273	
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273	
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273	
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273	

	ptratio	lstat	medv
0	15.3	4.98	24.0
1	17.8	9.14	21.6
2	17.8	4.03	34.7
3	18.7	2.94	33.4
4	18.7	5.33	36.2

```

..      ...      ...
501      21.0      9.67      22.4
502      21.0      9.08      20.6
503      21.0      5.64      23.9
504      21.0      6.48      22.0
505      21.0      7.88      11.9

```

[506 rows x 13 columns]

```
[41]: terms = Boston.columns.drop("medv")
terms
```

```
[41]: Index(['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax',
          'ptratio', 'lstat'],
          dtype='object')
```

```
[42]: X = MS(terms).fit_transform(Boston)
model = sm.OLS(y, X)
results = model.fit()
summarize(results)
```

```
[42]:
```

	coef	std err	t	P> t
intercept	41.6173	4.936	8.431	0.000
crim	-0.1214	0.033	-3.678	0.000
zn	0.0470	0.014	3.384	0.001
indus	0.0135	0.062	0.217	0.829
chas	2.8400	0.870	3.264	0.001
nox	-18.7580	3.851	-4.870	0.000
rm	3.6581	0.420	8.705	0.000
age	0.0036	0.013	0.271	0.787
dis	-1.4908	0.202	-7.394	0.000
rad	0.2894	0.067	4.325	0.000
tax	-0.0127	0.004	-3.337	0.001
ptratio	-0.9375	0.132	-7.091	0.000
lstat	-0.5520	0.051	-10.897	0.000

- Age has a high p-value. So how about we drop it from the predictors?

```
[43]: minus_age = Boston.columns.drop(["medv", "age"])
Xma = MS(minus_age).fit_transform(Boston)
model1 = sm.OLS(y, Xma)
summarize(model1.fit())
```

```
[43]:
```

	coef	std err	t	P> t
intercept	41.5251	4.920	8.441	0.000
crim	-0.1214	0.033	-3.683	0.000
zn	0.0465	0.014	3.379	0.001
indus	0.0135	0.062	0.217	0.829

chas	2.8528	0.868	3.287	0.001
nox	-18.4851	3.714	-4.978	0.000
rm	3.6811	0.411	8.951	0.000
dis	-1.5068	0.193	-7.825	0.000
rad	0.2879	0.067	4.322	0.000
tax	-0.0127	0.004	-3.333	0.001
ptratio	-0.9346	0.132	-7.099	0.000
lstat	-0.5474	0.048	-11.483	0.000

```
[44]: np.unique(Boston["indus"])
```

```
[44]: array([ 0.46,  0.74,  1.21,  1.22,  1.25,  1.32,  1.38,  1.47,  1.52,
          1.69,  1.76,  1.89,  1.91,  2.01,  2.02,  2.03,  2.18,  2.24,
          2.25,  2.31,  2.46,  2.68,  2.89,  2.93,  2.95,  2.97,  3.24,
          3.33,  3.37,  3.41,  3.44,  3.64,  3.75,  3.78,  3.97,  4.   ,
          4.05,  4.15,  4.39,  4.49,  4.86,  4.93,  4.95,  5.13,  5.19,
          5.32,  5.64,  5.86,  5.96,  6.06,  6.07,  6.09,  6.2  ,  6.41,
          6.91,  6.96,  7.07,  7.38,  7.87,  8.14,  8.56,  9.69,  9.9  ,
          10.01, 10.59, 10.81, 11.93, 12.83, 13.89, 13.92, 15.04, 18.1  ,
          19.58, 21.89, 25.65, 27.74])
```

Similarly, `indus` has a high p-value. Let's drop it as well. `minus_age_indus = Boston.columns.drop(["medv", "age", "indus"]) Xmai = MS(minus_age_indus).fit_transform(Boston) model1 = sm.OLS(y, Xmai) results1 = model1.fit() summarize(results1)`

We can also observe the F-statistic for the regression.

```
[45]: (results1.fvalue, results1.f_pvalue)
```

```
[45]: (308.9693351215988, 2.9820335524722154e-88)
```

1.9.16 Multivariate Goodness of Fit

1.9.17 We can access the individual components of results by name.

```
[46]: dir(results1)
```

```
[46]: ['HCO_se',
      'HC1_se',
      'HC2_se',
      'HC3_se',
      '_HCCM',
      '__class__',
      '__delattr__',
      '__dict__',
      '__dir__',
      '__doc__',
      '__eq__',
```

```

'__format__',
'__ge__',
'__getattribute__',
'__getstate__',
'__gt__',
'__hash__',
'__init__',
'__init_subclass__',
'__le__',
'__lt__',
'__module__',
'__ne__',
'__new__',
'__reduce__',
'__reduce_ex__',
'__repr__',
'__setattr__',
'__sizeof__',
'__str__',
'__subclasshook__',
'__weakref__',
'_abat_diagonal',
'_cache',
'_data_attr',
'_data_in_cache',
'_get_robustcov_results',
'_get_wald_nonlinear',
'_is_nested',
'_transform_predict_exog',
'_use_t',
'_wexog_singular_values',
'aic',
'bic',
'bse',
'centered_tss',
'compare_f_test',
'compare_lm_test',
'compare_lr_test',
'condition_number',
'conf_int',
'conf_int_el',
'cov_HC0',
'cov_HC1',
'cov_HC2',
'cov_HC3',
'cov_kwds',
'cov_params',

```

```
'cov_type',
'df_model',
'df_resid',
'diagn',
'eigenvals',
'el_test',
'ess',
'f_pvalue',
'f_test',
'fittedvalues',
'fvalue',
'get_influence',
'get_prediction',
'get_robustcov_results',
'info_criteria',
'initialize',
'k_constant',
'llf',
'load',
'model',
'mse_model',
'mse_resid',
'mse_total',
'nobs',
'normalized_cov_params',
'outlier_test',
'params',
'predict',
'pvalues',
'remove_data',
'resid',
'resid_pearson',
'rsquared',
'rsquared_adj',
'save',
'scale',
'ssr',
'summary',
'summary2',
't_test',
't_test_pairwise',
'tvalues',
'uncentered_tss',
'use_t',
'wald_test',
'wald_test_terms',
'wresid']
```

- `results.rsquared` gives us the R^2 and `np.sqrt(results.scale)` gives us the RSE.

```
[47]: print("RSE", np.sqrt(results1.scale))
```

```
RSE 6.173136281359115
```

```
[48]: ("R", results1.rsquared)
```

```
[48]: ('R', 0.5512689379421002)
```

- Variance Inflation Factors are sometimes useful to assess the collinearity effect in our regression model.

1.9.18 Compute VIFs and List Comprehension

```
[49]: vals = [VIF(X, i) for i in range(1, X.shape[1])]
      print(vals)
```

```
[1.7674859154310127, 2.2984589077358097, 3.9871806307570994, 1.071167773758404,
4.369092622844793, 1.9125324374368868, 3.0882320397311966, 3.954036641628298,
7.445300760069838, 9.002157663471797, 1.7970595931297786, 2.8707765008417514]
```

```
[50]: vif = pd.DataFrame({"vif": vals}, index=X.columns[1:])
      print(vif)
      ("VIF Range:", np.min(vif), np.max(vif))
```

	vif
crim	1.767486
zn	2.298459
indus	3.987181
chas	1.071168
nox	4.369093
rm	1.912532
age	3.088232
dis	3.954037
rad	7.445301
tax	9.002158
ptratio	1.797060
lstat	2.870777

```
[50]: ('VIF Range:', 1.071167773758404, 9.002157663471797)
```

- The VIFs are not very large.

1.9.19 Interaction terms

```
[51]: X = MS(["lstat", "age", ("lstat", "age")]).fit_transform(Boston)
      model2 = sm.OLS(y, X)
      results2 = model2.fit()
      summarize(results2)
```

```
[51]:
```

	coef	std err	t	P> t
intercept	36.0885	1.470	24.553	0.000
lstat	-1.3921	0.167	-8.313	0.000
age	-0.0007	0.020	-0.036	0.971
lstat:age	0.0042	0.002	2.244	0.025

```
[52]: (results2.rsquared, " > ", results1.rsquared)
```

```
[52]: (0.5557265450993936, ' > ', 0.5512689379421002)
```

- The interaction terms lstat:age are not statistically significant at 0.01 level of significance, and R2 does not significantly explain the variation in the model. Suffice to say, the interaction term can be dropped.

1.9.20 Non-linear transformation of the predictors

- The poly() function specifies the first argument term to be added to the model matrix

```
[53]: X = MS([poly("lstat", degree=2), "age"]).fit_transform(Boston)
model3 = sm.OLS(y, X)
results3 = model3.fit()
summarize(results3)
```

```
[53]:
```

	coef	std err	t	P> t
intercept	17.7151	0.781	22.681	0.0
poly(lstat, degree=2)[0]	-179.2279	6.733	-26.620	0.0
poly(lstat, degree=2)[1]	72.9908	5.482	13.315	0.0
age	0.0703	0.011	6.471	0.0

The effectively 0 p-value associated with the quadratic term suggests an improved model. The R2 confirms it

```
[54]: print(results3.rsquared, " > ", results2.rsquared)
```

```
0.6683791720749932 > 0.5557265450993936
```

- By default, poly() creates a basis matrix for inclusion in the model matrix whose columns are orthogonal polynomials which are designed for stable least squares computations. If we had included another argument, raw = True, the basis matrix would consist of lstat and lstat ** 2. Both represent quadratic polynomials. The fitted values would not change. Just the polynomial coefficients. The columns created by poly() do not include an intercept column. These are provided by MS().

1.9.21 Questions:

- What are orthogonal polynomials?
- <http://home.iitk.ac.in/~shalab/regression/Chapter12-Regression-PolynomialRegression.pdf>
- <https://stats.stackexchange.com/questions/258307/raw-or-orthogonal-polynomial-regression>

```
[55]: X = MS([poly("lstat", degree=2, raw=True), "age"]).fit_transform(Boston)
model3 = sm.OLS(y, X)
results3 = model3.fit()
summarize(results3)
```

```
[55]:
```

	coef	std err	t	P> t
intercept	41.2885	0.873	47.284	0.0
poly(lstat, degree=2, raw=True)[0]	-2.6883	0.131	-20.502	0.0
poly(lstat, degree=2, raw=True)[1]	0.0495	0.004	13.315	0.0
age	0.0703	0.011	6.471	0.0

```
[56]: print(results3.rsquared, " > ", results1.rsquared)
```

```
0.6683791720749932 > 0.5512689379421002
```

- Use the `anova_lm()` function to further quantify the superiority of the quadratic fit.

```
[57]: anova_lm(results1, results3)
```

```
[57]:
```


	df_resid	ssr	df_diff	ss_diff	F	Pr(>F)
0	503.0	19168.128609	0.0	NaN	NaN	NaN
1	502.0	14165.613251	1.0	5002.515357	177.278785	7.468491e-35

- `results1` corresponds to the linear model containing predictors `lstat` and `age` only.
- `results3` includes the quadratic term in `lstat`.
- The `anova_lm()` function performs a hypothesis test on the two models.
- H_0 : The quadratic term in the model is not needed.
- H_a : The larger model including the quadratic term is superior.
- Here, the F-statistic is 177.28 and the associated p-value is 0.
- The F-statistic is the t-statistic squared for the quadratic term in `results3`.
- These nested models differ by 1 degree of freedom.
- This provides very clear evidence that the quadratic term improves the model.
- The `anova_lm()` function can take more than two models as input.
- The comparison is successive pair-wise.
- That explains the NaNs in the first row of the output above, since there is no previous model with which to compare the output.

1.9.22 We can further plot the residuals of the regression against the fitted values to check if there still is a pattern discernible.

```
[58]: _, ax = subplots(figsize=(8, 8))
ax.scatter(results3.fittedvalues, results3.resid)
ax.set_xlabel("Fitted values")
ax.set_ylabel("Residuals")
ax.axhline(0, c="k", ls="--")
```

```
[58]: <matplotlib.lines.Line2D at 0x7370a0d620c0>
```

1.9.23 We can also try and add the interaction term (lstat, age) to the regression and check the results.

```
[59]: X = MS([poly("lstat", degree=2, raw=True), "age", ("lstat", "age"))].
      ↪fit_transform(
        Boston
      )
      model4 = sm.OLS(y, X)
      results4 = model4.fit()
      summarize(results4)
```

```
[59]:
```

	coef	std err	t	P> t
intercept	37.2658	1.250	29.816	0.0
poly(lstat, degree=2, raw=True)[0]	-2.2980	0.156	-14.723	0.0

```
poly(lstat, degree=2, raw=True)[1]    0.0584    0.004   14.015    0.0
age                                0.1439    0.020    7.279    0.0
lstat:age                         -0.0079    0.002   -4.424    0.0
```

```
[60]: print(results4.rsquared, " > ", results3.rsquared)
```

```
0.6808467217930462 > 0.6683791720749932
```

- The R2 in the interaction model again does not exceedingly explain the variance in the model compared to simply having the quadratic term.

1.9.24 Qualitative Predictors

1.10 Carseats data

```
[61]: Carseats = load_data("Carseats")
Carseats.columns
```

```
[61]: Index(['Sales', 'CompPrice', 'Income', 'Advertising', 'Population', 'Price',
          'ShelveLoc', 'Age', 'Education', 'Urban', 'US'],
          dtype='object')
```

```
[62]: Carseats.shape
```

```
[62]: (400, 11)
```

```
[63]: Carseats.describe()
```

```
[63]:
```

	Sales	CompPrice	Income	Advertising	Population	\
count	400.000000	400.000000	400.000000	400.000000	400.000000	
mean	7.496325	124.975000	68.657500	6.635000	264.840000	
std	2.824115	15.334512	27.986037	6.650364	147.376436	
min	0.000000	77.000000	21.000000	0.000000	10.000000	
25%	5.390000	115.000000	42.750000	0.000000	139.000000	
50%	7.490000	125.000000	69.000000	5.000000	272.000000	
75%	9.320000	135.000000	91.000000	12.000000	398.500000	
max	16.270000	175.000000	120.000000	29.000000	509.000000	

	Price	Age	Education
count	400.000000	400.000000	400.000000
mean	115.795000	53.322500	13.900000
std	23.676664	16.200297	2.620528
min	24.000000	25.000000	10.000000
25%	100.000000	39.750000	12.000000
50%	117.000000	54.500000	14.000000
75%	131.000000	66.000000	16.000000
max	191.000000	80.000000	18.000000

- ModelSpec() generates dummy variables for categorical columns automatically. This is termed a one-hot encoding of the categorical feature.

- Their columns sum to one. To avoid collinearity with the intercept, the first column is dropped.

1.10.1 Below we fit a multiple regression model with interaction terms.

```
[64]: allvars = list(Carseats.columns.drop("Sales"))
y = Carseats["Sales"]
final = allvars + [("Income", "Advertising"), ("Price", "Age")]
X = MS(final).fit_transform(Carseats)
model = sm.OLS(y, X)
summarize(model.fit())
```

```
[64]:
```

	coef	std err	t	P> t
intercept	6.5756	1.009	6.519	0.000
CompPrice	0.0929	0.004	22.567	0.000
Income	0.0109	0.003	4.183	0.000
Advertising	0.0702	0.023	3.107	0.002
Population	0.0002	0.000	0.433	0.665
Price	-0.1008	0.007	-13.549	0.000
ShelveLoc[Good]	4.8487	0.153	31.724	0.000
ShelveLoc[Medium]	1.9533	0.126	15.531	0.000
Age	-0.0579	0.016	-3.633	0.000
Education	-0.0209	0.020	-1.063	0.288
Urban[Yes]	0.1402	0.112	1.247	0.213
US[Yes]	-0.1576	0.149	-1.058	0.291
Income:Advertising	0.0008	0.000	2.698	0.007
Price:Age	0.0001	0.000	0.801	0.424

- It can be seen that ShelveLoc is significant and a good shelving location is associated with high sales (relative to a bad location). Medium has a smaller coefficient than Good leading us to believe that it leads to higher sales than a bad location, but lesser than a good location.

```
[65]: allDone()
```

```
<IPython.lib.display.Audio object>
```