

Lab: Linear Regression

Table of contents

| | |
|---|----|
| Set up IPython libraries for customizing notebook display | 2 |
| Import standard libraries | 2 |
| New imports | 2 |
| Import statsmodels objects | 2 |
| Import ISLP objects | 2 |
| Import User Functions | 3 |
| Inspecting objects and namespaces | 3 |
| Simple Linear Regression | 8 |
| We will use the Boston housing dataset which is in the package ISLP | 8 |
| Use sm.OLS to fit a simple linear regression | 8 |
| Extract the response and fit the model. | 9 |
| Summarize the results using the ISLP method summarize | 9 |
| Using Transformations: Fit and Transform | 9 |
| Full and exhaustive summary of the fit | 10 |
| Fitted coefficients can be retrieved as the <i>params</i> attribute of results | 10 |
| Computing predictions | 10 |
| We can predict confidence intervals for the predicted values. | 11 |
| We can obtain prediction intervals for the values which are wider than the confidence intervals since they're for a specific instance of lstat by setting obs=True. | 11 |
| Plot medv and lstat using DataFrame.plot.scatter() and add the regression line to the resulting plot. | 11 |
| Find the fitted values and residuals of the fit as attributes of the results object as <i>results.fittedvalues</i> and <i>results.resid</i> | 12 |
| Leverage statistics can be computed for any number of predictors using the <i>hat_matrix_diag</i> attribute of the value returned by the <i>get_influence()</i> method. | 13 |
| For a more conservative cutoff values for <i>hat_diag</i> , we have the following influence point(s): | 18 |
| Using DFFITS cutoff, we have the following influential points | 18 |
| Using Cooks Distance, we have the following influential points | 19 |
| Using Cooks Distance p-values, we have the following influential points | 19 |
| Using DFBeta for intercept, we have the following influential points | 19 |
| Using DFBeta for lstat, we have the following influential points | 20 |
| Multiple linear regression | 20 |
| Multivariate Goodness of Fit | 25 |

| | |
|--|----|
| We can access the individual components of results by name. | 25 |
| Compute VIFs and List Comprehension | 28 |
| Interaction terms | 28 |
| Non-linear transformation of the predictors | 29 |
| Questions: | 29 |
| We can further plot the residuals of the regression against the fitted values to check of there still is a pattern discernible. | 30 |
| We can also try and add the interaction term (lstat, age) to the regression and check the results. | 31 |
| Qualitative Predictors | 32 |
| Carseats data | 32 |
| Below we fit a multiple regression model with interaction terms. | 33 |

Set up IPython libraries for customizing notebook display

```
from notebookfuncs import *
```

Import standard libraries

```
import numpy as np
import pandas as pd
from matplotlib.pyplot import subplots
```

New imports

```
import statsmodels.api as sm
```

Import statsmodels objects

```
from statsmodels.stats.outliers_influence import variance_inflation_factor as VIF
from statsmodels.stats.anova import anova_lm
```

Import ISLP objects

```
from ISLP import load_data
from ISLP.models import ModelSpec as MS, summarize, poly
```

Import User Functions

```
from userfuncs import *
```

Inspecting objects and namespaces

```
dir()
```

```
['Audio',  
 'In',  
 'InteractiveShell',  
 'Latex',  
 'MS',  
 'Markdown',  
 'Math',  
 'Out',  
 'VIF',  
 '__builtin__',  
 '__builtins__',  
 '__name__',  
 '__spec__',  
 '_dh',  
 '_i',  
 '_i1',  
 '_i2',  
 '_i3',  
 '_i4',  
 '_i5',  
 '_i6',  
 '_i7',  
 '_ih',  
 '_ii',  
 '_iii',  
 '_oh',  
 'allDone',  
 'anova_lm',  
 'calculate_VIFs',  
 'check_symmetric',  
 'display',  
 'display_DFFITS_plot',  
 'display_cooks_distance_plot',  
 'display_hat_leverage_cutoffs',  
 'display_hat_leverage_plot',  
 'display_residuals_plot',
```

```

'display_studentized_residuals',
'dmatrices',
'exit',
'get_influence_points',
'get_ipython',
'get_results_df',
'identify_highest_VIF_feature',
'identify_least_significant_feature',
'influence_plot',
'is_numeric_dtype',
'is_pos_def',
'is_symmetric_pos_def',
'load_data',
'np',
'ojs_define',
'open',
'pd',
'perform_analysis',
'poly',
'printlatex',
'printmd',
'px',
'quit',
'sm',
'smf',
'standardize',
'stats',
'subplots',
'summarize']

```

```

A = np.array([3, 5, 11])
dir(A)

```

```

['T',
 '__abs__',
 '__add__',
 '__and__',
 '__array__',
 '__array_finalize__',
 '__array_function__',
 '__array_interface__',
 '__array_prepare__',
 '__array_priority__',
 '__array_struct__',
 '__array_ufunc__',
 '__array_wrap__',
 '__bool__',

```

```

'__buffer__',
'__class__',
'__class_getitem__',
'__complex__',
'__contains__',
'__copy__',
'__deepcopy__',
'__delattr__',
'__delitem__',
'__dir__',
'__divmod__',
'__dlpack__',
'__dlpack_device__',
'__doc__',
'__eq__',
'__float__',
'__floordiv__',
'__format__',
'__ge__',
'__getattr__',
'__getitem__',
'__getstate__',
'__gt__',
'__hash__',
'__iadd__',
'__iand__',
'__ifloordiv__',
'__ilshift__',
'__imatmul__',
'__imod__',
'__imul__',
'__index__',
'__init__',
'__init_subclass__',
'__int__',
'__invert__',
'__ior__',
'__ipow__',
'__irshift__',
'__isub__',
'__iter__',
'__itruediv__',
'__ixor__',
'__le__',
'__len__',
'__lshift__',

```

```
'__lt__',
'__matmul__',
'__mod__',
'__mul__',
'__ne__',
'__neg__',
'__new__',
'__or__',
'__pos__',
'__pow__',
'__radd__',
'__rand__',
'__rdivmod__',
'__reduce__',
'__reduce_ex__',
'__repr__',
'__rfloordiv__',
'__rlshift__',
'__rmatmul__',
'__rmod__',
'__rmul__',
'__ror__',
'__rpow__',
'__rrshift__',
'__rshift__',
'__rsub__',
'__rtruediv__',
'__rxor__',
'__setattr__',
'__setitem__',
'__setstate__',
'__sizeof__',
'__str__',
'__sub__',
'__subclasshook__',
'__truediv__',
'__xor__',
'all',
'any',
'argmax',
'argmin',
'argpartition',
'argsort',
'astype',
'base',
'byteswap',
```

'choose',
'clip',
'compress',
'conj',
'conjugate',
'copy',
'ctypes',
'cumprod',
'cumsum',
'data',
'diagonal',
'dot',
'dtype',
'dump',
'dumps',
'fill',
'flags',
'flat',
'flatten',
'getfield',
'imag',
'item',
'itemset',
'itemsizes',
'max',
'mean',
'min',
'nbytes',
'ndim',
'newbyteorder',
'nonzero',
'partition',
'prod',
'ptp',
'put',
'ravel',
'real',
'repeat',
'reshape',
'resize',
'round',
'searchsorted',
'setfield',
'setflags',
'shape',
'size',


```
'sort',
'squeeze',
'std',
'strides',
'sum',
'swapaxes',
'take',
'tobytes',
'tofile',
'tolist',
'tostring',
'trace',
'transpose',
'var',
'view']
```

```
A.sum()
```

```
19
```

Simple Linear Regression

We will use the Boston housing dataset which is in the package ISLP

```
Boston = load_data("Boston")
Boston.columns
```

```
Index(['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax',
      'ptratio', 'lstat', 'medv'],
      dtype='object')
```

```
len(Boston.columns)
```

```
13
```

```
# Boston?
```

Use sm.OLS to fit a simple linear regression

```
X = pd.DataFrame({"intercept": np.ones(Boston.shape[0]), "lstat": Boston["lstat"]})
X.head()
```

| | intercept | lstat |
|---|-----------|-------|
| 0 | 1.0 | 4.98 |
| 1 | 1.0 | 9.14 |

| | intercept | lstat |
|---|-----------|-------|
| 2 | 1.0 | 4.03 |
| 3 | 1.0 | 2.94 |
| 4 | 1.0 | 5.33 |

Extract the response and fit the model.

```
y = Boston["medv"]
model = sm.OLS(y, X)
results = model.fit()
```

<statsmodels.regression.linear_model.RegressionResultsWrapper at 0x7d9a3255e810>

Summarize the results using the ISLP method summarize

```
summarize(results)
```

| | coef | std err | t | P> t |
|-----------|---------|---------|---------|------|
| intercept | 34.5538 | 0.563 | 61.415 | 0.0 |
| lstat | -0.9500 | 0.039 | -24.528 | 0.0 |

Using Transformations: Fit and Transform

```
design = MS(["lstat"])
design = design.fit(Boston)
X = design.transform(Boston)
X.head()
```

| | intercept | lstat |
|---|-----------|-------|
| 0 | 1.0 | 4.98 |
| 1 | 1.0 | 9.14 |
| 2 | 1.0 | 4.03 |
| 3 | 1.0 | 2.94 |
| 4 | 1.0 | 5.33 |

```
design = MS(["lstat"])
design = design.fit_transform(Boston)
X.head()
```

| | intercept | lstat |
|---|-----------|-------|
| 0 | 1.0 | 4.98 |
| 1 | 1.0 | 9.14 |
| 2 | 1.0 | 4.03 |
| 3 | 1.0 | 2.94 |
| 4 | 1.0 | 5.33 |

Full and exhaustive summary of the fit

```
results.summary()
```

| | | | | | | |
|-------------------|------------------|---------------------|----------|-------|--------|--------|
| Dep. Variable: | medv | R-squared: | 0.544 | | | |
| Model: | OLS | Adj. R-squared: | 0.543 | | | |
| Method: | Least Squares | F-statistic: | 601.6 | | | |
| Date: | Tue, 25 Feb 2025 | Prob (F-statistic): | 5.08e-88 | | | |
| Time: | 14:38:50 | Log-Likelihood: | -1641.5 | | | |
| No. Observations: | 506 | AIC: | 3287. | | | |
| Df Residuals: | 504 | BIC: | 3295. | | | |
| Df Model: | 1 | | | | | |
| Covariance Type: | nonrobust | | | | | |
| | | | | | | |
| | coef | std err | t | P> t | [0.025 | 0.975] |
| intercept | 34.5538 | 0.563 | 61.415 | 0.000 | 33.448 | 35.659 |
| lstat | -0.9500 | 0.039 | -24.528 | 0.000 | -1.026 | -0.874 |
| | | | | | | |
| Omnibus: | 137.043 | Durbin-Watson: | 0.892 | | | |
| Prob(Omnibus): | 0.000 | Jarque-Bera (JB): | 291.373 | | | |
| Skew: | 1.453 | Prob(JB): | 5.36e-64 | | | |
| Kurtosis: | 5.319 | Cond. No. | 29.7 | | | |

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Fitted coefficients can be retrieved as the *params* attribute of results

```
results.params
```

```
intercept    34.553841
lstat        -0.950049
dtype: float64
```

Computing predictions

```
design = MS(["lstat"])
new_df = pd.DataFrame({"lstat": [5, 10, 15]})
```

```
print(new_df)
design = design.fit(new_df)
newX = design.transform(new_df)
newX
```

```
    lstat
0      5
1     10
2     15
```

| | intercept | lstat |
|---|-----------|-------|
| 0 | 1.0 | 5 |
| 1 | 1.0 | 10 |
| 2 | 1.0 | 15 |

```
new_predictions = results.get_prediction(newX)
new_predictions.predicted_mean
```

```
array([29.80359411, 25.05334734, 20.30310057])
```

We can predict confidence intervals for the predicted values.

```
new_predictions.conf_int(alpha=0.05)
```

```
array([[29.00741194, 30.59977628],
       [24.47413202, 25.63256267],
       [19.73158815, 20.87461299]])
```

We can obtain prediction intervals for the values which are wider than the confidence intervals since they're for a specific instance of lstat by setting `obs=True`.

```
new_predictions.conf_int(obs=True, alpha=0.05)
```

```
array([[17.56567478, 42.04151344],
       [12.82762635, 37.27906833],
       [ 8.0777421 , 32.52845905]])
```


Plot `medv` and `lstat` using `DataFrame.plot.scatter()` and add the regression line to the resulting plot.

```
ax = Boston.plot.scatter("lstat", "medv")
ax.axline(
```

```

(ax.get_xlim()[0], results.params.iloc[0]),
slope=results.params.iloc[1],
color="r",
linewidth=3,
)

```



- There is some evidence of non-linearity in the relationship b/w lstat and medv.


Find the fitted values and residuals of the fit as attributes of the results object as *results.fittedvalues* and *results.resid*.

- The `get_influence()` method computes various influence measures of the regression.

```

_, ax = subplots(figsize=(8, 8))
ax.scatter(results.fittedvalues, results.resid)
ax.set_xlabel("Fitted values")
ax.set_ylabel("Residuals")
ax.axhline(0, c="k", ls="--")


```



- On the basis of the residual plot, there is some evidence of non-linearity.

Leverage statistics can be computed for any number of predictors using the `hat_matrix_diag` attribute of the value returned by the `get_influence()` method.

```
display_hat_leverage_cutoffs(results)
```





```
display_hat_leverage_plot(results)
```


Unable to display output for mime type(s): text/html


Unable to display output for mime type(s): text/html

```
display_cooks_distance_plot(results)
```




```
display_DFFITS_plot(results)
```





```
inf_df, _ = get_influence_points(results)
inf_df
```

```
n = 506.0, p = 2
Average Hat Leverage: 0.003952569169960474
Hat Leverage Cutoff = 2 * Average Hat Leverage = 0.007905138339920948
DFBetas Cutoff = 3 / sqrt(n) = 0.1333662673423161
DFFITS Cutoff = 2 * sqrt(p/n) = 0.1257389226923863
Cooks Distance Cutoff = 1.0
Cooks Distance p-value Cutoff = 0.05
Studentized Residuals Cutoff = 3.0
Studentized Residuals p-value Cutoff = 0.01
```

| | dfb_intercept | dfb_lstat | cooks_d | hat_diag | student_resid | dffits | student_resid_pvalue | hat_influence |
|-----|---------------|-----------|----------|----------|---------------|----------|----------------------|---------------|
| 161 | 0.226022 | -0.189640 | 0.025315 | 0.006609 | 2.776857 | 0.226503 | 0.002847 | 0.018353 |
| 162 | 0.225506 | -0.188310 | 0.025219 | 0.006450 | 2.806416 | 0.226112 | 0.002602 | 0.018100 |
| 163 | 0.219862 | -0.176384 | 0.024252 | 0.005359 | 3.024654 | 0.222011 | 0.001308 | 0.016208 |
| 166 | 0.217770 | -0.172500 | 0.023921 | 0.005089 | 3.084034 | 0.220566 | 0.001077 | 0.015694 |
| 186 | 0.212939 | -0.164025 | 0.023201 | 0.004589 | 3.201426 | 0.217377 | 0.000727 | 0.014692 |
| 195 | 0.221577 | -0.179716 | 0.024534 | 0.005617 | 2.970018 | 0.223224 | 0.001560 | 0.016683 |
| 203 | 0.199749 | -0.157618 | 0.020220 | 0.005013 | 2.853126 | 0.202515 | 0.002254 | 0.014302 |
| 204 | 0.221985 | -0.180535 | 0.024602 | 0.005685 | 2.955977 | 0.223517 | 0.001632 | 0.016805 |

| | dfb_intercept | dfb_lstat | cooks_d | hat_diag | student_resid | dffits | student_resid_pvalue | hat_influence |
|-----|---------------|-----------|----------|----------|---------------|----------|----------------------|---------------|
| 214 | -0.197509 | 0.297576 | 0.051465 | 0.013063 | 2.807647 | 0.323011 | 0.002592 | 0.036676 |
| 225 | 0.211641 | -0.161830 | 0.023017 | 0.004476 | 3.229640 | 0.216554 | 0.000660 | 0.014455 |
| 228 | 0.178647 | -0.140414 | 0.016255 | 0.004938 | 2.573814 | 0.181309 | 0.005172 | 0.012709 |
| 233 | 0.196791 | -0.154507 | 0.019680 | 0.004918 | 2.841931 | 0.199782 | 0.002333 | 0.013975 |
| 257 | 0.207834 | -0.155541 | 0.022502 | 0.004180 | 3.306527 | 0.214221 | 0.000506 | 0.013821 |
| 261 | 0.128039 | -0.084194 | 0.009646 | 0.003106 | 2.501385 | 0.139616 | 0.006344 | 0.007769 |
| 262 | 0.189166 | -0.136023 | 0.019246 | 0.003742 | 3.231093 | 0.198020 | 0.000657 | 0.012090 |
| 267 | 0.184377 | -0.119477 | 0.020008 | 0.003032 | 3.672331 | 0.202505 | 0.000133 | 0.011133 |
| 280 | 0.164168 | -0.129770 | 0.013717 | 0.005047 | 2.335776 | 0.166365 | 0.009947 | 0.011789 |
| 283 | 0.220671 | -0.177936 | 0.024384 | 0.005476 | 2.999671 | 0.222580 | 0.001418 | 0.016425 |
| 368 | 0.220170 | -0.176971 | 0.024302 | 0.005402 | 3.015284 | 0.222227 | 0.001349 | 0.016290 |
| 369 | 0.217594 | -0.172182 | 0.023894 | 0.005068 | 3.088724 | 0.220447 | 0.001061 | 0.015654 |
| 370 | 0.221623 | -0.179808 | 0.024542 | 0.005625 | 2.968458 | 0.223257 | 0.001568 | 0.016697 |
| 371 | 0.155509 | -0.078029 | 0.018381 | 0.002355 | 4.004703 | 0.194572 | 0.000036 | 0.009431 |
| 372 | 0.165278 | -0.091836 | 0.018763 | 0.002529 | 3.901020 | 0.196431 | 0.000054 | 0.009866 |
| 374 | -0.294291 | 0.401657 | 0.086162 | 0.026865 | 2.511537 | 0.417300 | 0.006166 | 0.067473 |
| 412 | -0.253809 | 0.357605 | 0.070029 | 0.020290 | 2.615542 | 0.376405 | 0.004588 | 0.053070 |

For a more conservative cutoff values for `hat_diag`, we have the following influence point(s):

```
inf_df[inf_df["hat_diag"] > (3 * np.mean(inf_df["hat_diag"]))]
```

| | dfb_intercept | dfb_lstat | cooks_d | hat_diag | student_resid | dffits | student_resid_pvalue | hat_influence |
|-----|---------------|-----------|----------|----------|---------------|----------|----------------------|---------------|
| 374 | -0.294291 | 0.401657 | 0.086162 | 0.026865 | 2.511537 | 0.417300 | 0.006166 | 0.067473 |
| 412 | -0.253809 | 0.357605 | 0.070029 | 0.020290 | 2.615542 | 0.376405 | 0.004588 | 0.053070 |

Using DFFITS cutoff, we have the following influential points

```
inf_df[inf_df["dffits"] > 2 * np.sqrt(len(results.params) / results.nobs)]
```

| | dfb_intercept | dfb_lstat | cooks_d | hat_diag | student_resid | dffits | student_resid_pvalue | hat_influence |
|-----|---------------|-----------|----------|----------|---------------|----------|----------------------|---------------|
| 161 | 0.226022 | -0.189640 | 0.025315 | 0.006609 | 2.776857 | 0.226503 | 0.002847 | 0.018353 |
| 162 | 0.225506 | -0.188310 | 0.025219 | 0.006450 | 2.806416 | 0.226112 | 0.002602 | 0.018100 |
| 163 | 0.219862 | -0.176384 | 0.024252 | 0.005359 | 3.024654 | 0.222011 | 0.001308 | 0.016208 |
| 166 | 0.217770 | -0.172500 | 0.023921 | 0.005089 | 3.084034 | 0.220566 | 0.001077 | 0.015694 |
| 186 | 0.212939 | -0.164025 | 0.023201 | 0.004589 | 3.201426 | 0.217377 | 0.000727 | 0.014692 |
| 195 | 0.221577 | -0.179716 | 0.024534 | 0.005617 | 2.970018 | 0.223224 | 0.001560 | 0.016683 |
| 203 | 0.199749 | -0.157618 | 0.020220 | 0.005013 | 2.853126 | 0.202515 | 0.002254 | 0.014302 |
| 204 | 0.221985 | -0.180535 | 0.024602 | 0.005685 | 2.955977 | 0.223517 | 0.001632 | 0.016805 |

| | dfb_intercept | dfb_lstat | cooks_d | hat_diag | student_resid | dffits | student_resid_pvalue | hat_influence |
|-----|---------------|-----------|----------|----------|---------------|----------|----------------------|---------------|
| 214 | -0.197509 | 0.297576 | 0.051465 | 0.013063 | 2.807647 | 0.323011 | 0.002592 | 0.036676 |
| 225 | 0.211641 | -0.161830 | 0.023017 | 0.004476 | 3.229640 | 0.216554 | 0.000660 | 0.014455 |
| 228 | 0.178647 | -0.140414 | 0.016255 | 0.004938 | 2.573814 | 0.181309 | 0.005172 | 0.012709 |
| 233 | 0.196791 | -0.154507 | 0.019680 | 0.004918 | 2.841931 | 0.199782 | 0.002333 | 0.013975 |
| 257 | 0.207834 | -0.155541 | 0.022502 | 0.004180 | 3.306527 | 0.214221 | 0.000506 | 0.013821 |
| 261 | 0.128039 | -0.084194 | 0.009646 | 0.003106 | 2.501385 | 0.139616 | 0.006344 | 0.007769 |
| 262 | 0.189166 | -0.136023 | 0.019246 | 0.003742 | 3.231093 | 0.198020 | 0.000657 | 0.012090 |
| 267 | 0.184377 | -0.119477 | 0.020008 | 0.003032 | 3.672331 | 0.202505 | 0.000133 | 0.011133 |
| 280 | 0.164168 | -0.129770 | 0.013717 | 0.005047 | 2.335776 | 0.166365 | 0.009947 | 0.011789 |
| 283 | 0.220671 | -0.177936 | 0.024384 | 0.005476 | 2.999671 | 0.222580 | 0.001418 | 0.016425 |
| 368 | 0.220170 | -0.176971 | 0.024302 | 0.005402 | 3.015284 | 0.222227 | 0.001349 | 0.016290 |
| 369 | 0.217594 | -0.172182 | 0.023894 | 0.005068 | 3.088724 | 0.220447 | 0.001061 | 0.015654 |
| 370 | 0.221623 | -0.179808 | 0.024542 | 0.005625 | 2.968458 | 0.223257 | 0.001568 | 0.016697 |
| 371 | 0.155509 | -0.078029 | 0.018381 | 0.002355 | 4.004703 | 0.194572 | 0.000036 | 0.009431 |
| 372 | 0.165278 | -0.091836 | 0.018763 | 0.002529 | 3.901020 | 0.196431 | 0.000054 | 0.009866 |
| 374 | -0.294291 | 0.401657 | 0.086162 | 0.026865 | 2.511537 | 0.417300 | 0.006166 | 0.067473 |
| 412 | -0.253809 | 0.357605 | 0.070029 | 0.020290 | 2.615542 | 0.376405 | 0.004588 | 0.053070 |

Using Cooks Distance, we have the following influential points

```
inf_df[inf_df["cooks_d"] > 1.0]
```

| | dfb_intercept | dfb_lstat | cooks_d | hat_diag | student_resid | dffits | student_resid_pvalue | hat_influence |
|--|---------------|-----------|---------|----------|---------------|--------|----------------------|---------------|
|--|---------------|-----------|---------|----------|---------------|--------|----------------------|---------------|

Using Cooks Distance p-values, we have the following influential points

```
inf_df[inf_df["cooks_d_pvalue"] < 0.05]
```

| | dfb_intercept | dfb_lstat | cooks_d | hat_diag | student_resid | dffits | student_resid_pvalue | hat_influence |
|--|---------------|-----------|---------|----------|---------------|--------|----------------------|---------------|
|--|---------------|-----------|---------|----------|---------------|--------|----------------------|---------------|

Using DFBeta for intercept, we have the following influential points

```
inf_df[inf_df["dfb_intercept"] > (3 / np.sqrt(results.nobs))]
```

| | dfb_intercept | dfb_lstat | cooks_d | hat_diag | student_resid | dffits | student_resid_pvalue | hat_influence |
|-----|---------------|-----------|----------|----------|---------------|----------|----------------------|---------------|
| 161 | 0.226022 | -0.189640 | 0.025315 | 0.006609 | 2.776857 | 0.226503 | 0.002847 | 0.018353 |
| 162 | 0.225506 | -0.188310 | 0.025219 | 0.006450 | 2.806416 | 0.226112 | 0.002602 | 0.018100 |
| 163 | 0.219862 | -0.176384 | 0.024252 | 0.005359 | 3.024654 | 0.222011 | 0.001308 | 0.016208 |

| | dfb_intercept | dfb_lstat | cooks_d | hat_diag | student_resid | dffits | student_resid_pvalue | hat_influence |
|-----|---------------|-----------|----------|----------|---------------|----------|----------------------|---------------|
| 166 | 0.217770 | -0.172500 | 0.023921 | 0.005089 | 3.084034 | 0.220566 | 0.001077 | 0.015694 |
| 186 | 0.212939 | -0.164025 | 0.023201 | 0.004589 | 3.201426 | 0.217377 | 0.000727 | 0.014692 |
| 195 | 0.221577 | -0.179716 | 0.024534 | 0.005617 | 2.970018 | 0.223224 | 0.001560 | 0.016683 |
| 203 | 0.199749 | -0.157618 | 0.020220 | 0.005013 | 2.853126 | 0.202515 | 0.002254 | 0.014302 |
| 204 | 0.221985 | -0.180535 | 0.024602 | 0.005685 | 2.955977 | 0.223517 | 0.001632 | 0.016805 |
| 225 | 0.211641 | -0.161830 | 0.023017 | 0.004476 | 3.229640 | 0.216554 | 0.000660 | 0.014455 |
| 228 | 0.178647 | -0.140414 | 0.016255 | 0.004938 | 2.573814 | 0.181309 | 0.005172 | 0.012709 |
| 233 | 0.196791 | -0.154507 | 0.019680 | 0.004918 | 2.841931 | 0.199782 | 0.002333 | 0.013975 |
| 257 | 0.207834 | -0.155541 | 0.022502 | 0.004180 | 3.306527 | 0.214221 | 0.000506 | 0.013821 |
| 262 | 0.189166 | -0.136023 | 0.019246 | 0.003742 | 3.231093 | 0.198020 | 0.000657 | 0.012090 |
| 267 | 0.184377 | -0.119477 | 0.020008 | 0.003032 | 3.672331 | 0.202505 | 0.000133 | 0.011133 |
| 280 | 0.164168 | -0.129770 | 0.013717 | 0.005047 | 2.335776 | 0.166365 | 0.009947 | 0.011789 |
| 283 | 0.220671 | -0.177936 | 0.024384 | 0.005476 | 2.999671 | 0.222580 | 0.001418 | 0.016425 |
| 368 | 0.220170 | -0.176971 | 0.024302 | 0.005402 | 3.015284 | 0.222227 | 0.001349 | 0.016290 |
| 369 | 0.217594 | -0.172182 | 0.023894 | 0.005068 | 3.088724 | 0.220447 | 0.001061 | 0.015654 |
| 370 | 0.221623 | -0.179808 | 0.024542 | 0.005625 | 2.968458 | 0.223257 | 0.001568 | 0.016697 |
| 371 | 0.155509 | -0.078029 | 0.018381 | 0.002355 | 4.004703 | 0.194572 | 0.000036 | 0.009431 |
| 372 | 0.165278 | -0.091836 | 0.018763 | 0.002529 | 3.901020 | 0.196431 | 0.000054 | 0.009866 |

Using DFBeta for lstat, we have the following influential points


```
inf_df[inf_df["dfb_lstat"] > (3 / np.sqrt(results.nobs))]
```

| | dfb_intercept | dfb_lstat | cooks_d | hat_diag | student_resid | dffits | student_resid_pvalue | hat_influence |
|-----|---------------|-----------|----------|----------|---------------|----------|----------------------|---------------|
| 214 | -0.197509 | 0.297576 | 0.051465 | 0.013063 | 2.807647 | 0.323011 | 0.002592 | 0.036676 |
| 374 | -0.294291 | 0.401657 | 0.086162 | 0.026865 | 2.511537 | 0.417300 | 0.006166 | 0.067473 |
| 412 | -0.253809 | 0.357605 | 0.070029 | 0.020290 | 2.615542 | 0.376405 | 0.004588 | 0.053070 |

Multiple linear regression


```
Boston.plot.scatter("age", "medv")
X = MS(["lstat", "age"]).fit_transform(Boston)
model1 = sm.OLS(y, X)
results1 = model1.fit()
summarize(results1)
```

| | coef | std err | t | P> t |
|-----------|---------|---------|---------|-------|
| intercept | 33.2228 | 0.731 | 45.458 | 0.000 |
| lstat | -1.0321 | 0.048 | -21.416 | 0.000 |
| age | 0.0345 | 0.012 | 2.826 | 0.005 |




```
Boston["logage"] = np.log(Boston["age"])
Boston.plot.scatter("logage", "medv")
X = MS(["lstat", "logage"]).fit_transform(Boston)
model1 = sm.OLS(y, X)
resultslog = model1.fit()
print(summarize(resultslog))
```

| | coef | std err | t | P> t |
|-----------|---------|---------|---------|------|
| intercept | 30.2143 | 1.947 | 15.517 | 0.00 |
| lstat | -1.0051 | 0.045 | -22.213 | 0.00 |
| logage | 1.2312 | 0.529 | 2.327 | 0.02 |



```
Boston["sqrtage"] = np.sqrt(Boston["age"])
Boston.plot.scatter("sqrtage", "medv")
X = MS(["lstat", "sqrtage"]).fit_transform(Boston)
model1 = sm.OLS(y, X)
resultssqrt = model1.fit()
summarize(resultssqrt)
```

| | coef | std err | t | P> t |
|-----------|---------|---------|---------|-------|
| intercept | 31.8635 | 1.174 | 27.139 | 0.000 |
| lstat | -1.0203 | 0.047 | -21.703 | 0.000 |
| sqrtage | 0.4450 | 0.171 | 2.606 | 0.009 |



```
Boston = Boston.drop(columns=["logage", "sqrtage"])
```

| | crim | zn | indus | chas | nox | rm | age | dis | rad | tax | ptratio | lstat | medv |
|-----|---------|------|-------|------|-------|-------|------|--------|-----|-----|---------|-------|------|
| 0 | 0.00632 | 18.0 | 2.31 | 0 | 0.538 | 6.575 | 65.2 | 4.0900 | 1 | 296 | 15.3 | 4.98 | 24.0 |
| 1 | 0.02731 | 0.0 | 7.07 | 0 | 0.469 | 6.421 | 78.9 | 4.9671 | 2 | 242 | 17.8 | 9.14 | 21.6 |
| 2 | 0.02729 | 0.0 | 7.07 | 0 | 0.469 | 7.185 | 61.1 | 4.9671 | 2 | 242 | 17.8 | 4.03 | 34.7 |
| 3 | 0.03237 | 0.0 | 2.18 | 0 | 0.458 | 6.998 | 45.8 | 6.0622 | 3 | 222 | 18.7 | 2.94 | 33.4 |
| 4 | 0.06905 | 0.0 | 2.18 | 0 | 0.458 | 7.147 | 54.2 | 6.0622 | 3 | 222 | 18.7 | 5.33 | 36.2 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 501 | 0.06263 | 0.0 | 11.93 | 0 | 0.573 | 6.593 | 69.1 | 2.4786 | 1 | 273 | 21.0 | 9.67 | 22.4 |
| 502 | 0.04527 | 0.0 | 11.93 | 0 | 0.573 | 6.120 | 76.7 | 2.2875 | 1 | 273 | 21.0 | 9.08 | 20.6 |
| 503 | 0.06076 | 0.0 | 11.93 | 0 | 0.573 | 6.976 | 91.0 | 2.1675 | 1 | 273 | 21.0 | 5.64 | 23.9 |
| 504 | 0.10959 | 0.0 | 11.93 | 0 | 0.573 | 6.794 | 89.3 | 2.3889 | 1 | 273 | 21.0 | 6.48 | 22.0 |
| 505 | 0.04741 | 0.0 | 11.93 | 0 | 0.573 | 6.030 | 80.8 | 2.5050 | 1 | 273 | 21.0 | 7.88 | 11.9 |

```
terms = Boston.columns.drop("medv")
terms
```

```
Index(['crim', 'zn', 'indus', 'chas', 'nox', 'rm', 'age', 'dis', 'rad', 'tax',
      'ptratio', 'lstat'],
      dtype='object')
```

```
X = MS(terms).fit_transform(Boston)
model = sm.OLS(y, X)
```



```
results = model.fit()
summarize(results)
```

| | coef | std err | t | P> t |
|-----------|----------|---------|---------|-------|
| intercept | 41.6173 | 4.936 | 8.431 | 0.000 |
| crim | -0.1214 | 0.033 | -3.678 | 0.000 |
| zn | 0.0470 | 0.014 | 3.384 | 0.001 |
| indus | 0.0135 | 0.062 | 0.217 | 0.829 |
| chas | 2.8400 | 0.870 | 3.264 | 0.001 |
| nox | -18.7580 | 3.851 | -4.870 | 0.000 |
| rm | 3.6581 | 0.420 | 8.705 | 0.000 |
| age | 0.0036 | 0.013 | 0.271 | 0.787 |
| dis | -1.4908 | 0.202 | -7.394 | 0.000 |
| rad | 0.2894 | 0.067 | 4.325 | 0.000 |
| tax | -0.0127 | 0.004 | -3.337 | 0.001 |
| ptratio | -0.9375 | 0.132 | -7.091 | 0.000 |
| lstat | -0.5520 | 0.051 | -10.897 | 0.000 |

- Age has a high p-value. So how about we drop it from the predictors?

```
minus_age = Boston.columns.drop(["medv", "age"])
Xma = MS(minus_age).fit_transform(Boston)
model1 = sm.OLS(y, Xma)
summarize(model1.fit())
```

| | coef | std err | t | P> t |
|-----------|----------|---------|---------|-------|
| intercept | 41.5251 | 4.920 | 8.441 | 0.000 |
| crim | -0.1214 | 0.033 | -3.683 | 0.000 |
| zn | 0.0465 | 0.014 | 3.379 | 0.001 |
| indus | 0.0135 | 0.062 | 0.217 | 0.829 |
| chas | 2.8528 | 0.868 | 3.287 | 0.001 |
| nox | -18.4851 | 3.714 | -4.978 | 0.000 |
| rm | 3.6811 | 0.411 | 8.951 | 0.000 |
| dis | -1.5068 | 0.193 | -7.825 | 0.000 |
| rad | 0.2879 | 0.067 | 4.322 | 0.000 |
| tax | -0.0127 | 0.004 | -3.333 | 0.001 |
| ptratio | -0.9346 | 0.132 | -7.099 | 0.000 |
| lstat | -0.5474 | 0.048 | -11.483 | 0.000 |

```
np.unique(Boston["indus"])
```

```
array([ 0.46,  0.74,  1.21,  1.22,  1.25,  1.32,  1.38,  1.47,  1.52,
        1.69,  1.76,  1.89,  1.91,  2.01,  2.02,  2.03,  2.18,  2.24,
```

```

2.25, 2.31, 2.46, 2.68, 2.89, 2.93, 2.95, 2.97, 3.24,
3.33, 3.37, 3.41, 3.44, 3.64, 3.75, 3.78, 3.97, 4. ,
4.05, 4.15, 4.39, 4.49, 4.86, 4.93, 4.95, 5.13, 5.19,
5.32, 5.64, 5.86, 5.96, 6.06, 6.07, 6.09, 6.2 , 6.41,
6.91, 6.96, 7.07, 7.38, 7.87, 8.14, 8.56, 9.69, 9.9 ,
10.01, 10.59, 10.81, 11.93, 12.83, 13.89, 13.92, 15.04, 18.1 ,
19.58, 21.89, 25.65, 27.74])

```

Similarly, indus has a high p-value. Let's drop it as well.

```

minus_age_indus = Boston.columns.drop(["medv", "age", "indus"]) Xmai = MS(minus_age_indus).fit_transform(Boston)
model1 = sm.OLS(y, Xmai) results1 = model1.fit() summarize(results1)

```

We can also observe the F-statistic for the regression.

```

(results1.fvalue, results1.f_pvalue)

```

```

(308.9693351215988, 2.9820335524722154e-88)

```

Multivariate Goodness of Fit

We can access the individual components of results by name.

```

dir(results1)

```

```

['HCO_se',
 'HC1_se',
 'HC2_se',
 'HC3_se',
 '_HCCM',
 '__class__',
 '__delattr__',
 '__dict__',
 '__dir__',
 '__doc__',
 '__eq__',
 '__format__',
 '__ge__',
 '__getattr__',
 '__getstate__',
 '__gt__',
 '__hash__',
 '__init__',
 '__init_subclass__',
 '__le__',
 '__lt__',
 '__module__',
 '__ne__',

```

```

'__new__',
'__reduce__',
'__reduce_ex__',
'__repr__',
'__setattr__',
'__sizeof__',
'__str__',
'__subclasshook__',
'__weakref__',
'_abat_diagonal',
'_cache',
'_data_attr',
'_data_in_cache',
'_get_robustcov_results',
'_get_wald_nonlinear',
'_is_nested',
'_transform_predict_exog',
'_use_t',
'_wexog_singular_values',
'aic',
'bic',
'bse',
'centered_tss',
'compare_f_test',
'compare_lm_test',
'compare_lr_test',
'condition_number',
'conf_int',
'conf_int_el',
'cov_HC0',
'cov_HC1',
'cov_HC2',
'cov_HC3',
'cov_kwds',
'cov_params',
'cov_type',
'df_model',
'df_resid',
'diagn',
'eigenvals',
'el_test',
'ess',
'f_pvalue',
'f_test',
'fittedvalues',
'fvalue',

```

```

'get_influence',
'get_prediction',
'get_robustcov_results',
'info_criteria',
'initialize',
'k_constant',
'llf',
'load',
'model',
'mse_model',
'mse_resid',
'mse_total',
'nobs',
'normalized_cov_params',
'outlier_test',
'params',
'predict',
'pvalues',
'remove_data',
'resid',
'resid_pearson',
'rsquared',
'rsquared_adj',
'save',
'scale',
'ssr',
'summary',
'summary2',
't_test',
't_test_pairwise',
'tvalues',
'uncentered_tss',
'use_t',
'wald_test',
'wald_test_terms',
'wresid']

```

- results.rsquared gives us the R2 and np.sqrt(results.scale) gives us the RSE.

```
print("RSE", np.sqrt(results1.scale))
```

```
RSE 6.173136281359115
```

```
("R", results1.rsquared)
```

```
('R', 0.5512689379421002)
```

- Variance Inflation Factors are sometimes useful to assess the collinearity effect in our regression

model.

Compute VIFs and List Comprehension

```
vals = [VIF(X, i) for i in range(1, X.shape[1])]
print(vals)
```

```
[1.7674859154310127, 2.2984589077358097, 3.9871806307570994, 1.071167773758404, 4.369092622844793, 1.9125320125320125, 3.0882320125320125, 3.9540370125320125, 7.4453010125320125, 9.0021580125320125, 1.7970601253201253, 2.8707770125320125]
```

```
vif = pd.DataFrame({"vif": vals}, index=X.columns[1:])
print(vif)
("VIF Range:", np.min(vif), np.max(vif))
```

| | vif |
|---------|----------|
| crim | 1.767486 |
| zn | 2.298459 |
| indus | 3.987181 |
| chas | 1.071168 |
| nox | 4.369093 |
| rm | 1.912532 |
| age | 3.088232 |
| dis | 3.954037 |
| rad | 7.445301 |
| tax | 9.002158 |
| ptratio | 1.797060 |
| lstat | 2.870777 |

```
('VIF Range:', 1.071167773758404, 9.002157663471797)
```

- The VIFs are not very large.

Interaction terms

```
X = MS(["lstat", "age", ("lstat", "age")]).fit_transform(Boston)
model2 = sm.OLS(y, X)
results2 = model2.fit()
summarize(results2)
```

| | coef | std err | t | P> t |
|-----------|---------|---------|--------|-------|
| intercept | 36.0885 | 1.470 | 24.553 | 0.000 |
| lstat | -1.3921 | 0.167 | -8.313 | 0.000 |
| age | -0.0007 | 0.020 | -0.036 | 0.971 |
| lstat:age | 0.0042 | 0.002 | 2.244 | 0.025 |

```
(results2.rsquared, " > ", results1.rsquared)
```

```
(0.5557265450993936, ' > ', 0.5512689379421002)
```

- The interaction terms `lstat:age` are not statistically significant at 0.01 level of significance, and `R2` does not significantly explain the variation in the model. Suffice to say, the interaction term can be dropped.

Non-linear transformation of the predictors

- The `poly()` function specifies the first argument term to be added to the model matrix

```
X = MS([poly("lstat", degree=2), "age"]).fit_transform(Boston)
model3 = sm.OLS(y, X)
results3 = model3.fit()
summarize(results3)
```

| | coef | std err | t | P> t |
|--------------------------|-----------|---------|---------|------|
| intercept | 17.7151 | 0.781 | 22.681 | 0.0 |
| poly(lstat, degree=2)[0] | -179.2279 | 6.733 | -26.620 | 0.0 |
| poly(lstat, degree=2)[1] | 72.9908 | 5.482 | 13.315 | 0.0 |
| age | 0.0703 | 0.011 | 6.471 | 0.0 |

The effectively 0 p-value associated with the quadratic term suggests an improved model. The `R2` confirms it

```
print(results3.rsquared, " > ", results2.rsquared)
```

```
0.6683791720749932 > 0.5557265450993936
```

- By default, `poly()` creates a basis matrix for inclusion in the model matrix whose columns are orthogonal polynomials which are designed for stable least squares computations. If we had included another argument, `raw = True`, the basis matrix would consist of `lstat` and `lstat ** 2`. Both represent quadratic polynomials. The fitted values would not change. Just the polynomial coefficients. The columns created by `poly()` do not include an intercept column. These are provided by `MS()`.

Questions:

- What are orthogonal polynomials?
- <http://home.iitk.ac.in/~shalab/regression/Chapter12-Regression-PolynomialRegression.pdf>
- <https://stats.stackexchange.com/questions/258307/raw-or-orthogonal-polynomial-regression>

```
X = MS([poly("lstat", degree=2, raw=True), "age"]).fit_transform(Boston)
model3 = sm.OLS(y, X)
```

```
results3 = model3.fit()
summarize(results3)
```

| | coef | std err | t | P> t |
|------------------------------------|---------|---------|---------|------|
| intercept | 41.2885 | 0.873 | 47.284 | 0.0 |
| poly(lstat, degree=2, raw=True)[0] | -2.6883 | 0.131 | -20.502 | 0.0 |
| poly(lstat, degree=2, raw=True)[1] | 0.0495 | 0.004 | 13.315 | 0.0 |
| age | 0.0703 | 0.011 | 6.471 | 0.0 |

```
print(results3.rsquared, " > ", results1.rsquared)
```

0.6683791720749932 > 0.5512689379421002

- Use the `anova_lm()` function to further quantify the superiority of the quadratic fit.


```
anova_lm(results1, results3)
```

| | df_resid | ssr | df_diff | ss_diff | F | Pr(>F) |
|---|----------|--------------|---------|-------------|------------|--------------|
| 0 | 503.0 | 19168.128609 | 0.0 | NaN | NaN | NaN |
| 1 | 502.0 | 14165.613251 | 1.0 | 5002.515357 | 177.278785 | 7.468491e-35 |

- results1 corresponds to the linear model containing predictors lstat and age only.
- results3 includes the quadratic term in lstat.
- The `anova_lm()` function performs a hypothesis test on the two models.
- H0: The quadratic term in the model is not needed.
- Ha: The larger model including the quadratic term is superior.
- Here, the F-statistic is 177.28 and the associated p-value is 0.
- The F-statistic is the t-statistic squared for the quadratic term in results3.
- These nested models differ by 1 degree of freedom.
- This provides very clear evidence that the quadratic term improves the model.
- The `anova_lm()` function can take more than two models as input.
- The comparison is successive pair-wise.
- That explains the NaNs in the first row of the output above, since there is no previous model with which to compare the output.

We can further plot the residuals of the regression against the fitted values to check if there still is a pattern discernible.

```
_, ax = subplots(figsize=(8, 8))
ax.scatter(results3.fittedvalues, results3.resid)
ax.set_xlabel("Fitted values")
ax.set_ylabel("Residuals")
ax.axhline(0, c="k", ls="--")
```



We can also try and add the interaction term (lstat, age) to the regression and check the results.

```
X = MS([poly("lstat", degree=2, raw=True), "age", ("lstat", "age"))].fit_transform(
    Boston
)
model4 = sm.OLS(y, X)
results4 = model4.fit()
summarize(results4)
```


| | coef | std err | t | P> t |
|------------------------------------|---------|---------|---------|------|
| intercept | 37.2658 | 1.250 | 29.816 | 0.0 |
| poly(lstat, degree=2, raw=True)[0] | -2.2980 | 0.156 | -14.723 | 0.0 |
| poly(lstat, degree=2, raw=True)[1] | 0.0584 | 0.004 | 14.015 | 0.0 |
| age | 0.1439 | 0.020 | 7.279 | 0.0 |
| lstat:age | -0.0079 | 0.002 | -4.424 | 0.0 |

```
print(results4.rsquared, " > ", results3.rsquared)
```

```
0.6808467217930462 > 0.6683791720749932
```

- The R2 in the interaction model again does not exceedingly explain the variance in the model compared to simply having the quadratic term.

Qualitative Predictors

Carseats data

```
Carseats = load_data("Carseats")
Carseats.columns
```

```
Index(['Sales', 'CompPrice', 'Income', 'Advertising', 'Population', 'Price',
      'ShelveLoc', 'Age', 'Education', 'Urban', 'US'],
      dtype='object')
```

```
Carseats.shape
```

```
(400, 11)
```

```
Carseats.describe()
```

| | Sales | CompPrice | Income | Advertising | Population | Price | Age | Education |
|-------|------------|------------|------------|-------------|------------|------------|------------|------------|
| count | 400.000000 | 400.000000 | 400.000000 | 400.000000 | 400.000000 | 400.000000 | 400.000000 | 400.000000 |
| mean | 7.496325 | 124.975000 | 68.657500 | 6.635000 | 264.840000 | 115.795000 | 53.322500 | 13.900000 |
| std | 2.824115 | 15.334512 | 27.986037 | 6.650364 | 147.376436 | 23.676664 | 16.200297 | 2.620528 |
| min | 0.000000 | 77.000000 | 21.000000 | 0.000000 | 10.000000 | 24.000000 | 25.000000 | 10.000000 |
| 25% | 5.390000 | 115.000000 | 42.750000 | 0.000000 | 139.000000 | 100.000000 | 39.750000 | 12.000000 |
| 50% | 7.490000 | 125.000000 | 69.000000 | 5.000000 | 272.000000 | 117.000000 | 54.500000 | 14.000000 |
| 75% | 9.320000 | 135.000000 | 91.000000 | 12.000000 | 398.500000 | 131.000000 | 66.000000 | 16.000000 |
| max | 16.270000 | 175.000000 | 120.000000 | 29.000000 | 509.000000 | 191.000000 | 80.000000 | 18.000000 |

- ModelSpec() generates dummy variables for categorical columns automatically. This is termed a one-hot encoding of the categorical feature.

- Their columns sum to one. To avoid collinearity with the intercept, the first column is dropped.

Below we fit a multiple regression model with interaction terms.

```
allvars = list(Carseats.columns.drop("Sales"))
y = Carseats["Sales"]
final = allvars + [("Income", "Advertising"), ("Price", "Age")]
X = MS(final).fit_transform(Carseats)
model = sm.OLS(y, X)
summarize(model.fit())
```

| | coef | std err | t | P> t |
|--------------------|---------|---------|---------|-------|
| intercept | 6.5756 | 1.009 | 6.519 | 0.000 |
| CompPrice | 0.0929 | 0.004 | 22.567 | 0.000 |
| Income | 0.0109 | 0.003 | 4.183 | 0.000 |
| Advertising | 0.0702 | 0.023 | 3.107 | 0.002 |
| Population | 0.0002 | 0.000 | 0.433 | 0.665 |
| Price | -0.1008 | 0.007 | -13.549 | 0.000 |
| ShelveLoc[Good] | 4.8487 | 0.153 | 31.724 | 0.000 |
| ShelveLoc[Medium] | 1.9533 | 0.126 | 15.531 | 0.000 |
| Age | -0.0579 | 0.016 | -3.633 | 0.000 |
| Education | -0.0209 | 0.020 | -1.063 | 0.288 |
| Urban[Yes] | 0.1402 | 0.112 | 1.247 | 0.213 |
| US[Yes] | -0.1576 | 0.149 | -1.058 | 0.291 |
| Income:Advertising | 0.0008 | 0.000 | 2.698 | 0.007 |
| Price:Age | 0.0001 | 0.000 | 0.801 | 0.424 |

- It can be seen that ShelvLoc is significant and a good shelving location is associated with high sales (relative to a bad location). Medium has a smaller coefficient than Good leading us to believe that it leads to higher sales than a bad location, but lesser than a good location.

```
allDone()
```

```
<IPython.lib.display.Audio object>
```