

Marketing

February 21, 2025

1 Marketing Exercise

1.1 Import notebook functions

```
[1]: from notebookfuncs import *
```

1.2 Import standard libraries

```
[2]: import numpy as np
import pandas as pd
from matplotlib.pyplot import subplots
```

1.3 New imports

```
[3]: import statsmodels.api as sm
```

1.4 Import statsmodel.objects

```
[4]: from statsmodels.stats.outliers_influence import variance_inflation_factor as v
      ↪ VIF
from statsmodels.stats.outliers_influence import summary_table
from statsmodels.stats.anova import anova_lm
```

1.5 Import ISLP objects

```
[5]: import ISLP
from ISLP import models
from ISLP import load_data
from ISLP.models import ModelSpec as MS, summarize, poly
```

Inspecting objects and namespaces

```
[6]: dir()
```

```
[6]: ['Audio',
      'ISLP',
      'In',
      'InteractiveShell',
```

```
'Latex',
'MS',
'Markdown',
'Math',
'Out',
'VIF',
'_',
'__',
'___',
'__builtin__',
'__builtins__',
'__doc__',
'__loader__',
'__name__',
'__package__',
'__spec__',
'_dh',
'_i',
'_i1',
'_i2',
'_i3',
'_i4',
'_i5',
'_i6',
'_ih',
'_ii',
'_iii',
'_oh',
'allDone',
'anova_lm',
'display',
'exit',
'get_ipython',
'load_data',
'models',
'np',
'open',
'pd',
'poly',
'printlatex',
'printmd',
'quit',
'sm',
'subplots',
'summarize',
'summary_table']
```

```
[7]: Advertising = pd.read_csv("Advertising.csv")
# Drop first column
Advertising = Advertising.iloc[:, 1:]
Advertising.head()
```

```
[7]:      TV  Radio  Newspaper  Sales
0  230.1   37.8      69.2    22.1
1   44.5   39.3      45.1    10.4
2   17.2   45.9      69.3     9.3
3  151.5   41.3      58.5    18.5
4  180.8   10.8      58.4    12.9
```

```
[8]: Advertising.describe()
```

```
[8]:      TV      Radio  Newspaper      Sales
count  200.000000  200.000000  200.000000  200.000000
mean    147.042500   23.264000   30.554000   14.022500
std     85.854236   14.846809   21.778621    5.217457
min      0.700000    0.000000    0.300000    1.600000
25%     74.375000    9.975000   12.750000   10.375000
50%    149.750000   22.900000   25.750000   12.900000
75%    218.825000   36.525000   45.100000   17.400000
max    296.400000   49.600000  114.000000   27.000000
```

1.6 Is there a relationship between sales and advertising budget?

```
[9]: y = Advertising["Sales"]
cols = list(Advertising.columns)
cols.remove("Sales")
X = MS(cols).fit_transform(Advertising)
model = sm.OLS(y, X)
results = model.fit()
print("F-value", results.fvalue)
print("F-pvalue", results.f_pvalue)
summarize(results)
```

F-value 570.2707036590944

F-pvalue 1.575227256092416e-96

```
[9]:      coef  std err      t  P>|t|
intercept  2.9389    0.312   9.422   0.00
TV          0.0458    0.001  32.809   0.00
Radio       0.1885    0.009  21.893   0.00
Newspaper -0.0010    0.006  -0.177   0.86
```

```
[10]: dir(models)
```

```
[10]: ['Column',
      'Feature',
      'FeatureSelector',
      'ModelSpec',
      'Stepwise',
      'StringIO',
      '__builtins__',
      '__cached__',
      '__doc__',
      '__file__',
      '__loader__',
      '__name__',
      '__package__',
      '__path__',
      '__spec__',
      'bs',
      'build_columns',
      'columns',
      'contrast',
      'derived_feature',
      'generic_selector',
      'min_max_strategy',
      'model_spec',
      'np',
      'ns',
      'pca',
      'pd',
      'poly',
      'sklearn_selected',
      'sklearn_selection_path',
      'sklearn_sm',
      'sklearn_wrap',
      'strategy',
      'summarize']
```

- The p-value corresponding to the F-statistic is very low. Thus, clear evidence of a relationship between sales and advertising budget.

```
[11]: dir(results)
```

```
[11]: ['HCO_se',
      'HC1_se',
      'HC2_se',
      'HC3_se',
      '_HCCM',
      '__class__',
      '__delattr__',
      '__dict__',
```

```

'__dir__',
'__doc__',
'__eq__',
'__format__',
'__ge__',
'__getattribute__',
'__getstate__',
'__gt__',
'__hash__',
'__init__',
'__init_subclass__',
'__le__',
'__lt__',
'__module__',
'__ne__',
'__new__',
'__reduce__',
'__reduce_ex__',
'__repr__',
'__setattr__',
'__sizeof__',
'__str__',
'__subclasshook__',
'__weakref__',
'_abat_diagonal',
'_cache',
'_data_attr',
'_data_in_cache',
'_get_robustcov_results',
'_get_wald_nonlinear',
'_is_nested',
'_transform_predict_exog',
'_use_t',
'_wexog_singular_values',
'aic',
'bic',
'bse',
'centered_tss',
'compare_f_test',
'compare_lm_test',
'compare_lr_test',
'condition_number',
'conf_int',
'conf_int_el',
'cov_HC0',
'cov_HC1',
'cov_HC2',

```

'cov_HC3',
'cov_kwds',
'cov_params',
'cov_type',
'df_model',
'df_resid',
'diagn',
'eigenvals',
'el_test',
'ess',
'f_pvalue',
'f_test',
'fittedvalues',
'fvalue',
'get_influence',
'get_prediction',
'get_robustcov_results',
'info_criteria',
'initialize',
'k_constant',
'llf',
'load',
'model',
'mse_model',
'mse_resid',
'mse_total',
'nobs',
'normalized_cov_params',
'outlier_test',
'params',
'predict',
'pvalues',
'remove_data',
'resid',
'resid_pearson',
'rsquared',
'rsquared_adj',
'save',
'scale',
'ssr',
'summary',
'summary2',
't_test',
't_test_pairwise',
'tvalues',
'uncentered_tss',
'use_t',

```
'wald_test',
'wald_test_terms',
'wresid']
```

1.7 How strong is the relationship?

```
[12]: results.summary()
```

```
[12]:
```

Dep. Variable:	Sales	R-squared:	0.897
Model:	OLS	Adj. R-squared:	0.896
Method:	Least Squares	F-statistic:	570.3
Date:	Fri, 21 Feb 2025	Prob (F-statistic):	1.58e-96
Time:	19:19:32	Log-Likelihood:	-386.18
No. Observations:	200	AIC:	780.4
Df Residuals:	196	BIC:	793.6
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
intercept	2.9389	0.312	9.422	0.000	2.324	3.554
TV	0.0458	0.001	32.809	0.000	0.043	0.049
Radio	0.1885	0.009	21.893	0.000	0.172	0.206
Newspaper	-0.0010	0.006	-0.177	0.860	-0.013	0.011

Omnibus:	60.414	Durbin-Watson:	2.084
Prob(Omnibus):	0.000	Jarque-Bera (JB):	151.241
Skew:	-1.327	Prob(JB):	1.44e-33
Kurtosis:	6.332	Cond. No.	454.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[13]: y.mean()
```

```
[13]: 14.0225
```

```
[14]: results.resid.std()
```

```
[14]: 1.6727572743844117
```

```
[15]: (results.resid.std() / y.mean()) * 100
```

```
[15]: 11.929094486606608
```

- The residual standard error (RSE) is 1.67 and the mean value of the response is 14.023 which translates to a percentage error of roughly 11.93%

```
[16]: ("R-squared", results.rsquared, "Adjusted R-squared", results.rsquared_adj)
```

```
[16]: ('R-squared', 0.8972106381789522, 'Adjusted R-squared', 0.8956373316204668)
```

- The R2 explains about 90% of the variance in Sales.

1.8 Which media are associated with Sales?

- The low p-values for Radio and TV suggest that only they are related to Sales.

1.9 How large is the association between each medium and sales?

```
[17]: results.conf_int(alpha=0.05)
```

```
[17]:
```

	0	1
intercept	2.323762	3.554016
TV	0.043014	0.048516
Radio	0.171547	0.205513
Newspaper	-0.012616	0.010541

- The confidence intervals for TV and Radio are narrow and far from zero. This provides evidence that these media are related to sales.
- The interval for Newspaper includes zero indicating that it is not statistically significant given values of TV and Radio.

```
[18]: vals = [VIF(X, i) for i in range(1, X.shape[1])]
print(vals)
```

```
[1.00461078493965, 1.1449519171055353, 1.1451873787239288]
```

- The VIF scores are 1.005, 1.145 and 1.145 respectively for TV, radio and newspaper. These suggest no evidence of collinearity as an explanation for wide standard errors for newspaper.
- In order to assess the association of each medium individually on sales, we can perform three separate linear regressions.

```
[19]: TV = MS(["TV"]).fit_transform(Advertising)
model = sm.OLS(y, TV)
results = model.fit()
print(summarize(results))
Radio = MS(["Radio"]).fit_transform(Advertising)
model = sm.OLS(y, Radio)
results = model.fit()
print(summarize(results))
Newspaper = MS(["Newspaper"]).fit_transform(Advertising)
model = sm.OLS(y, Newspaper)
results = model.fit()
print(summarize(results))
```

	coef	std err	t	P> t
intercept	7.0326	0.458	15.360	0.0
TV	0.0475	0.003	17.668	0.0

	coef	std err	t	P> t
intercept	9.3116	0.563	16.542	0.0

Radio	0.2025	0.020	9.921	0.0
	coef	std err	t	P> t
intercept	12.3514	0.621	19.876	0.000
Newspaper	0.0547	0.017	3.300	0.001

Looking at the p-values, there is evidence of a strong association b/w TV and sales and radio and sales. There is evidence of a mild association between Newspaper and sales when TV and radio are ignored.

1.10 How accurately can we predict future sales?

- Given that \$100,000 is spent on TV advertising, and \$20,000 is spent on Radio advertising, we need to compute the 95% Confidence intervals for each city (i.e., the mean) and the prediction interval for a particular city (also at 95% confidence intervals).

1.10.1 Fit the regression dropping the Newspaper column as insignificant

```
[20]: y = Advertising["Sales"]
cols = list(Advertising.columns)
cols.remove("Sales")
cols.remove("Newspaper")
X = MS(cols).fit_transform(Advertising)
model = sm.OLS(y, X)
results = model.fit()
print("F-value", results.fvalue)
print("F-pvalue", results.f_pvalue)
summarize(results)
```

F-value 859.6177183058211

F-pvalue 4.8273618513354486e-98

```
[20]:      coef  std err      t  P>|t|
intercept  2.9211    0.294   9.919   0.0
TV          0.0458    0.001  32.909   0.0
Radio      0.1880    0.008  23.382   0.0
```

```
[21]: results.summary()
```

```
[21]:
```

Dep. Variable:	Sales	R-squared:	0.897
Model:	OLS	Adj. R-squared:	0.896
Method:	Least Squares	F-statistic:	859.6
Date:	Fri, 21 Feb 2025	Prob (F-statistic):	4.83e-98
Time:	19:19:33	Log-Likelihood:	-386.20
No. Observations:	200	AIC:	778.4
Df Residuals:	197	BIC:	788.3
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
intercept	2.9211	0.294	9.919	0.000	2.340	3.502
TV	0.0458	0.001	32.909	0.000	0.043	0.048
Radio	0.1880	0.008	23.382	0.000	0.172	0.204
Omnibus:		60.022	Durbin-Watson:		2.081	
Prob(Omnibus):		0.000	Jarque-Bera (JB):		148.679	
Skew:		-1.323	Prob(JB):		5.19e-33	
Kurtosis:		6.292	Cond. No.		425.	

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[22]: design = MS(["TV", "Radio"])
new_df = pd.DataFrame({"TV": [100], "Radio": [20]})
print(new_df)
new_X = design.fit_transform(new_df)
new_predictions = results.get_prediction(new_X)
new_predictions.predicted_mean
```

```
      TV  Radio
0  100    20
```

```
[22]: array([11.25646595])
```

1.10.2 We predict the confidence intervals at 95% as follows:

```
[23]: new_predictions.conf_int(alpha=0.05)
```

```
[23]: array([[10.98525445, 11.52767746]])
```

1.10.3 We predict the prediction interval for a particular city as follows:

```
[24]: new_predictions.conf_int(alpha=0.05, obs=True)
```


```
[24]: array([[ 7.92961607, 14.58331584]])
```

- Both intervals are centered at 11,256 but the prediction intervals are wider reflecting the additional uncertainty around sales for a particular city as against the average sales for many locations.

1.10.4 Is the relationship linear?


```
[25]: _, ax = subplots(figsize=(8, 8))
ax.scatter(results.fittedvalues, results.resid)
ax.set_xlabel("Fitted values")
ax.set_ylabel("Residuals")
ax.axhline(0, c="k", ls="--")
```

[25]: <matplotlib.lines.Line2D at 0x7bd0434ad790>




```
[26]: _, ax = subplots(figsize=(8, 8))
      ax.scatter(Advertising["TV"], results.resid)
      ax.set_xlabel("TV")
      ax.set_ylabel("Residuals")
      ax.axhline(0, c="k", ls="--")
```

[26]: <matplotlib.lines.Line2D at 0x7bd040ba93d0>



```
[27]: _, ax = subplots(figsize=(8, 8))
      ax.scatter(Advertising["Radio"], results.resid)
      ax.set_xlabel("Radio")
      ax.set_ylabel("Residuals")
      ax.axhline(0, c="k", ls="--")
```

```
[27]: <matplotlib.lines.Line2D at 0x7bd0434aee10>
```



- There is evidence of non-linearity in the model from the residuals plotted against the fitted values. Looking at the residuals versus predictors plots, it appears that TV is a better candidate for quadratification.

```
[28]: X = MS([poly("TV", degree=2, raw=True), "Radio"]).fit_transform(Advertising)
model = sm.OLS(y, X)
results = model.fit()
summarize(results)
```

```
[28]:
```

	coef	std err	t	P> t
intercept	1.2876	0.359000	3.588	0.0
poly(TV, degree=2, raw=True)[0]	0.0784	0.005000	15.736	0.0
poly(TV, degree=2, raw=True)[1]	-0.0001	0.000017	-6.775	0.0
Radio	0.1930	0.007000	26.465	0.0

```
[29]: results.summary()
```

```
[29]:
```

Dep. Variable:	Sales	R-squared:	0.917
Model:	OLS	Adj. R-squared:	0.915
Method:	Least Squares	F-statistic:	719.0
Date:	Fri, 21 Feb 2025	Prob (F-statistic):	1.80e-105
Time:	19:19:33	Log-Likelihood:	-365.16
No. Observations:	200	AIC:	738.3
Df Residuals:	196	BIC:	751.5
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
intercept	1.2876	0.359	3.588	0.000	0.580	1.995
poly(TV, degree=2, raw=True)[0]	0.0784	0.005	15.736	0.000	0.069	0.088
poly(TV, degree=2, raw=True)[1]	-0.0001	1.68e-05	-6.775	0.000	-0.000	-8.05e-05
Radio	0.1930	0.007	26.465	0.000	0.179	0.207

Omnibus:	19.524	Durbin-Watson:	2.136
Prob(Omnibus):	0.000	Jarque-Bera (JB):	44.712
Skew:	-0.413	Prob(JB):	1.95e-10
Kurtosis:	5.164	Cond. No.	1.29e+05


Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.29e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
[30]: _, ax = subplots(figsize=(8, 8))
ax.scatter(results.fittedvalues, results.resid)
ax.set_xlabel("Fitted values")
ax.set_ylabel("Residuals")
ax.axhline(0, c="k", ls="--")
```

```
[30]: <matplotlib.lines.Line2D at 0x7bd04098eba0>
```



While the fit has improved as seen from the R^2 increasing by 2 percentage points, there is still some non-linearity visible in the residuals plot against fitted values.

References:

https://www.kellogg.northwestern.edu/faculty/weber/emp/_session_3/nonlinearities.htm

<https://online.stat.psu.edu/stat462/node/120/>

1.10.5 Is there synergy among the advertising media?

Synergy implies an interaction effect. That's what we test out now.

```
[31]: X = MS([poly("TV", raw=True, degree=2), "Radio", ("TV", "Radio"))].
      ↪fit_transform(
        Advertising
      )
      model = sm.OLS(y, X)
      results = model.fit()
      summarize(results)
```

```
[31]:
```

	coef	std err	t	P> t
intercept	5.1371	0.193000	26.663	0.0
poly(TV, degree=2, raw=True)[0]	0.0509	0.002000	22.810	0.0
poly(TV, degree=2, raw=True)[1]	-0.0001	0.000007	-15.920	0.0
Radio	0.0352	0.006000	5.959	0.0
TV:Radio	0.0011	0.000035	31.061	0.0

```
[32]: results.summary()
```

```
[32]:
```

Dep. Variable:	Sales	R-squared:	0.986
Model:	OLS	Adj. R-squared:	0.986
Method:	Least Squares	F-statistic:	3432.
Date:	Fri, 21 Feb 2025	Prob (F-statistic):	1.79e-179
Time:	19:19:34	Log-Likelihood:	-186.86
No. Observations:	200	AIC:	383.7
Df Residuals:	195	BIC:	400.2
Df Model:	4		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
intercept	5.1371	0.193	26.663	0.000	4.757	5.517
poly(TV, degree=2, raw=True)[0]	0.0509	0.002	22.810	0.000	0.047	0.055
poly(TV, degree=2, raw=True)[1]	-0.0001	6.89e-06	-15.920	0.000	-0.000	-9.61e-05
Radio	0.0352	0.006	5.959	0.000	0.024	0.047
TV:Radio	0.0011	3.47e-05	31.061	0.000	0.001	0.001

Omnibus:	169.759	Durbin-Watson:	2.204
Prob(Omnibus):	0.000	Jarque-Bera (JB):	4031.167
Skew:	-2.988	Prob(JB):	0.00
Kurtosis:	24.166	Cond. No.	1.70e+05

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.7e+05. This might indicate that there are strong multicollinearity or other numerical problems.

- Finally, when we add an interaction term TV * Radio to the model, we can see that the residual fit exhibits no pattern. And the R2 is 98.6%.

1.10.6 Compute VIFs and List Comprehension

```
[33]: vals = [VIF(X, i) for i in range(1, X.shape[1])]
      print(vals)
```

```
[18.787830609925035, 15.885268501061871, 3.9253174186837008, 6.940088238444382]
```

```
[34]: vif = pd.DataFrame({"vif": vals}, index=X.columns[1:])
      print(vif)
      ("VIF Range:", np.min(vif), np.max(vif))
```

```

                                vif
poly(TV, degree=2, raw=True)[0]  18.787831
poly(TV, degree=2, raw=True)[1]  15.885269
Radio                           3.925317
TV:Radio                        6.940088
```

```
[34]: ('VIF Range:', 3.9253174186837008, 18.787830609925035)
```

- The VIF ranges are high. These can be reduced by transforming variables to mean 0.
- <https://stats.stackexchange.com/questions/23538/quadratic-term-and-variance-inflation-factor>

```
[35]: Advertising["TV"] = Advertising["TV"] - Advertising["TV"].mean()
      Advertising["Radio"] = Advertising["Radio"] - Advertising["Radio"].mean()
```

```
[36]: X = MS([poly("TV", raw=True, degree=2), "Radio", ("TV", "Radio"))].
      ↪fit_transform(
          Advertising
      )
      model = sm.OLS(y, X)
      results = model.fit()
      summarize(results)
```

```
[36]:
```

	coef	std err	t	P> t
intercept	14.7525	0.067000	219.634	0.0
poly(TV, degree=2, raw=True)[0]	0.0437	0.001000	84.111	0.0
poly(TV, degree=2, raw=True)[1]	-0.0001	0.000007	-15.920	0.0
Radio	0.1935	0.003000	64.526	0.0
TV:Radio	0.0011	0.000035	31.061	0.0

```
[37]: results.summary()
```

```
[37]:
```

Dep. Variable:	Sales	R-squared:	0.986
Model:	OLS	Adj. R-squared:	0.986
Method:	Least Squares	F-statistic:	3432.
Date:	Fri, 21 Feb 2025	Prob (F-statistic):	1.79e-179
Time:	19:19:34	Log-Likelihood:	-186.86
No. Observations:	200	AIC:	383.7
Df Residuals:	195	BIC:	400.2
Df Model:	4		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
intercept	14.7525	0.067	219.634	0.000	14.620	14.885
poly(TV, degree=2, raw=True)[0]	0.0437	0.001	84.111	0.000	0.043	0.045
poly(TV, degree=2, raw=True)[1]	-0.0001	6.89e-06	-15.920	0.000	-0.000	-9.61e-05
Radio	0.1935	0.003	64.526	0.000	0.188	0.199
TV:Radio	0.0011	3.47e-05	31.061	0.000	0.001	0.001

Omnibus:	169.759	Durbin-Watson:	2.204
Prob(Omnibus):	0.000	Jarque-Bera (JB):	4031.167
Skew:	-2.988	Prob(JB):	0.00
Kurtosis:	24.166	Cond. No.	1.49e+04

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.49e+04. This might indicate that there are strong multicollinearity or other numerical problems.

```
[38]: vals = [VIF(X, i) for i in range(1, X.shape[1])]
      print(vals)
```

```
[1.0172717815970211, 1.017084612216564, 1.013513326764562, 1.0075840215785734]
```

```
[39]: vif = pd.DataFrame({"vif": vals}, index=X.columns[1:])
      print(vif)
      ("VIF Range:", np.min(vif), np.max(vif))
```

```

                                vif
poly(TV, degree=2, raw=True)[0]  1.017272
poly(TV, degree=2, raw=True)[1]  1.017085
Radio                           1.013513
TV:Radio                        1.007584
```

```
[39]: ('VIF Range:', 1.0075840215785734, 1.0172717815970211)
```

```
[40]: allDone()
```

```
<IPython.lib.display.Audio object>
```